

# Ch8. Processes and Programs - Studying sh

Prof. Seokin Hong

Kyungpook National University

Fall 2018

# Objectives

---

## ■ Ideas and Skills

- What a Unix **shell** does
- The Unix model of a **process**
- How to **run a program**
- How to **create a process**
- How **parent and child processes communicate**

## ■ System Calls

- fork, exec, wait, exit

## ■ Commands

- sh, ps

# Contents

---

- 8.1 Process
- 8.2 Learning about Processes with ps
- 8.3 The Shell:  
A Tool for Process and Program Control
- 8.4 How the Shell Runs Programs
- 8.5 Writing a Shell: psh2.c

# Program and Process

---

## ■ Program


- A sequence of machine-language instructions stored in a file

## ■ Running a program means

- Load that list of machine-language instructions into memory
- The processor(CPU) execute the instructions one by one

## ■ Process = Programs in action

- *A process is an **instance of a computer program that is being executed**. It is a piece of program that is loaded, being executing and running in the system.*



# Learning about Processes with “ps” command

---

- **ps** : list current processes

```
[seokin@compasslab1:~$ ps
  PID TTY          TIME CMD
 23467 pts/0        00:00:00 bash
 23476 pts/0        00:00:00 ps
```

- **PID**
  - process ID
- **TTY**
  - the name of the console or terminal that the user logged into
- **TIME**
  - the amount of CPU time in minutes and seconds that the process has been running
- **CMD**
  - the name of the command that launched the process

# Learning about Processes with “ps” command

---

## ■ ps -a

```
[seokin@compasslab1:~$ ps -a
  PID TTY          TIME CMD
  973 tty1        00:00:00 gnome-session-b
  980 tty1        00:01:39 gnome-shell
  999 tty1        00:00:00 Xwayland
 1373 tty1        00:00:00 ibus-daemon
 1376 tty1        00:00:00 ibus-dconf
 1379 tty1        00:00:00 ibus-x11
 1416 tty1        00:00:00 gsd-xsettings
 1424 tty1        00:00:00 gsd-a11y-settin
 1432 tty1        00:00:00 gsd-clipboard
 1433 tty1        00:04:04 gsd-color
 1436 tty1        00:00:00 gsd-datetime
 1438 tty1        00:00:00 gsd-housekeepin
 1440 tty1        00:00:00 gsd-keyboard
 1441 tty1        00:00:00 gsd-media-keys
 1444 tty1        00:00:00 gsd-mouse
 1446 tty1        00:00:01 gsd-power
      .
      .
      .
      .
      .
      .
```

## -a option

- lists more processes, including ones being run by other users and at other terminals
- the output from `—a` does not include the shell

# Learning about Processes with “ps” command

---

## ■ ps -al

```
$ ps -la
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
000	S	504	1779	1731	0	69	0	-	1086	do_sel	pts/0	00:00:13	gv
000	S	504	1780	1779	0	69	0	-	2309	do_sel	pts/0	00:00:07	gs
000	S	504	1781	1731	0	72	0	-	1320	do_sel	pts/0	00:00:01	vi
000	S	519	2013	1993	0	69	19	-	1300	do_sel	pts/2	00:00:23	xpain
000	S	519	2017	1993	0	69	0	-	363	read_c	pts/2	00:00:02	mail
000	R	500	2023	1755	0	79	0	-	750	-	pts/1	00:00:00	ps

- **S** : status of each process.
  - S: sleeping, R : running
- **UID** : user ID
- **PID** : process ID
- **PPID** : parent process ID
- **PRI** : priority
  - higher numbers mean lower priority
- **NI** : niceness level
  - 19(nicest) ~ -20(not nice to other)
  - a process with a higher nice number will yield CPU time to other processes on the system.
  - ※ The kernel uses these values, PRI & NI, to decide when to run processes.

# Learning about Processes with “ps” command

---

## ■ ps -al

```
$ ps -la
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
000	S	504	1779	1731	0	69	0	-	1086	do_sel	pts/0	00:00:13	gv
000	S	504	1780	1779	0	69	0	-	2309	do_sel	pts/0	00:00:07	gs
000	S	504	1781	1731	0	72	0	-	1320	do_sel	pts/0	00:00:01	vi
000	S	519	2013	1993	0	69	19	-	1300	do_sel	pts/2	00:00:23	xpain
000	S	519	2017	1993	0	69	0	-	363	read_c	pts/2	00:00:02	mail
000	R	500	2023	1755	0	79	0	-	750	-	pts/1	00:00:00	ps

- **SZ** : process size (KB)
  - the amount of memory the process is using
- **WCHAN** : wait channel
  - Name of the kernel function in which the process is sleeping (waiting on)
  - All the processes in this example are waiting for input; the entries `read_c` and `do_sel` refer to the kernel function
- **ADDR and F**
  - no longer used, but appear in the output for compatibility with programs that expect to see them.



# Learning about Processes with “ps” command

---

## ■ ps -af

```
[seokin@compasslab1:~$ ps -af
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
gdm	973	969	0	10月 24	tty1	00:00:00	/usr/lib/gnome-session/gnome-session-b
gdm	980	973	0	10月 24	tty1	00:01:39	/usr/bin/gnome-shell
gdm	999	980	0	10月 24	tty1	00:00:00	/usr/bin/Xwayland :1024 -rootless -ter
gdm	1373	980	0	10月 24	tty1	00:00:00	ibus-daemon --xim --panel disable
gdm	1376	1373	0	10月 24	tty1	00:00:00	/usr/lib/ibus/ibus-dconf
gdm	1379	1	0	10月 24	tty1	00:00:00	/usr/lib/ibus/ibus-x11 --kill-daemon
gdm	1416	973	0	10月 24	tty1	00:00:00	/usr/lib/gnome-settings-daemon/gsd-xse
gdm	1424	973	0	10月 24	tty1	00:00:00	/usr/lib/gnome-settings-daemon/gsd-a11


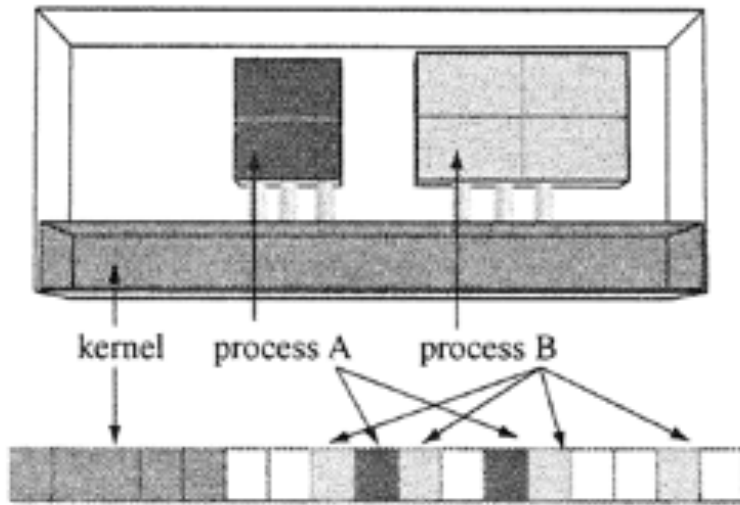
.....

```
seokin 23530 23467 0 23:36 pts/0 00:00:00 ps -af
```

- The username is displayed instead of the UID number
- Completed command line is listed in the CMD column

# Computer Memory & Computer Program

---



- Memory in a Unix system is divided into kernel space and user space
- Processes live in user space
- Processes are usually divided into smaller chunks, just as disk files are divided into disk blocks
- As a file has an allocation list of disk blocks, a process has a structure to hold the allocation list of memory pages

# Computer Memory & Computer Program

---

- Creating a process is similar to creating a disk file.
  - Kernel has to find some free pages of memory to hold the machine-language codes and data bytes for the program.
  - Kernel sets up some data structures to store memory allocation information and the attributes of the process.

# Contents

---

- 8.1 Process
- 8.2 Learning about Processes with `ps`
- 8.3 The Shell: A Tool for Process and Program Control
- 8.4 How the Shell Runs Programs
- 8.5 Writing a Shell: `psh2.c`

# THE SHELL

---

- A shell is a program that manages processes and runs programs
- Thee main functions of shells
  - (a) Shells run programs
  - (b) Shells manage input and output
  - (c) Shells can be programmed

# THE SHELL

---

```
[seokin@compasslab1:~$ ls
'[2018-2nd-sysp-001]HW1CopiaDirectory'
'[2018-2nd-sysp-001]HW1CopiaDirectory.zip'
a
add_user.py
a.out
Desktop
[seokin@compasslab1:~$ ps
  PID TTY          TIME CMD
 24041 pts/0    00:00:00 bash
 24050 pts/0    00:00:00 ps
```

dir1  
Documents  
Downloads  
examples.desktop  
hw1  
lab\_dir\_creation

log  
Music  
Pictures  
play\_again1.c  
Public  
student\_id

system\_programming  
Templates  
test  
Videos  
workspace

## ■ Running Programs

- The shell loads programs into memory and runs them
- Shell = program launcher

# THE SHELL

---

```
[seokin@compasslab1:~$ ls
'[2018-2nd-sysp-001]HW1CopaDirectory'  dir1      log        student_id
'[2018-2nd-sysp-001]HW1CopaDirectory.zip' Documents  ls.txt     system_programming
a                                       Downloads Music      Templates
add_user.py                           examples.desktop Pictures    test
a.out                                  hw1        play_again1.c Videos
Desktop                               lab_dir_creation Public      workspace

[seokin@compasslab1:~$ ls > ls.txt
[seokin@compasslab1:~$ ls | grep P
Pictures
Public
seokin@compasslab1:~$ █
```

## ■ Managing Input and Output

- `>`, `<`, `|` are symbols for input/output redirection
- With the symbols, a user tells the shell to attach the input and output of processes to disk file or to other processes

# THE SHELL

---

```
[seokin@compasslab1:~$ for ((num =0; num <3 ; num++))  
[> do  
[> echo "num:$num"  
[> done  
num:0  
num:1  
num:2
```

## ■ Programming

- The shell is also a programming language with variables and flow control (if, while, etc.)
- We can make an executable script (shell script)



# Contents

---

- 8.1 Process
- 8.2 Learning about Processes with ps
- 8.3 The Shell: A Tool for Process and Program Control
- 8.4 How the Shell Runs Programs
- 8.5 Writing a Shell: psh2.c

# How the shell runs programs

---




FIGURE 8.4

A user asks a shell to run a program.

- A. The user types a.out
- B. The shell creates a new process to run the program
- C. The shell loads the program from the disk into the process
- D. The program runs in its process until it is done

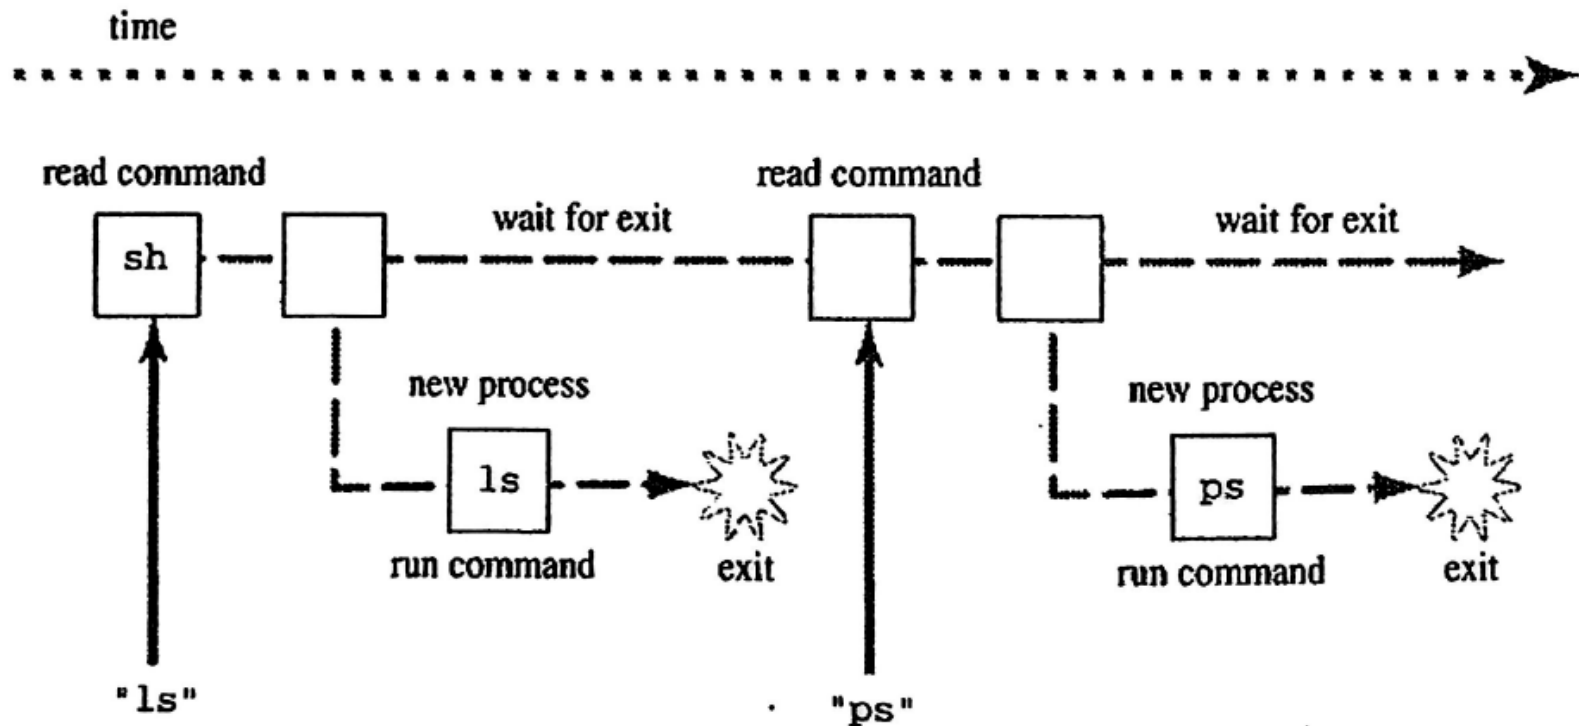
# The Main Loop of a Shell

- Shell consists of following loop:

```
while ( ! end_of_input )  
    get command  
    execute command  
    wait for command to finish
```

To write a shell, we need to learn how to

1. Run a program
2. Create a process
3. Wait for exit()



# Q1: How Does a Program Run a Program?

---

- Answer: The program calls “execvp()’ system call

---

## execvp

PURPOSE	Execute a file, with PATH searching
---------	-------------------------------------

INCLUDE	#include <unistd.h>
---------	---------------------

USAGE	result = execvp(const char *file, const char *argv[])
-------	---

ARGS	file    name of file to execute argv    array of strings
------	---

RETURNS	-1    if error
---------	----------------


# Q1: How Does a Program Run a Program?

---

How Unix runs programs:

```
execvp(progname, arglist)
```

1. copies the named program into the calling process,
2. passes the specified list of strings to the program as `argv[]`, then
3. runs the program.



1. Program calls `execvp`
2. Kernel loads program from disk into the process
3. Kernel copies `arglist` into the process
4. Kernel calls `main(argc, argv)`

# Example 1

---

```
/* execl.c - shows how easy it is for a program to run a program
*/
#include <stdio.h>
#include <unistd.h>

main()
{
    char    *arglist[3];
    arglist[0] = "ls";           //First string is the name of the program
    arglist[1] = "-l";
    arglist[2] = 0 ;            //Last string must have a null pointer
    printf("* * * About to exec ls -l\n");
    execvp( "ls" , arglist );
    printf("* * * ls is done. bye\n");
}
```

# Example 1

---

```
*** About to exec 'ls -l'
total 60
drwxrwxr-x 2 seokin seokin 4096  9月 12 01:50 lab1
drwxrwxr-x 2 seokin seokin 4096  9月 12 01:50 lab10
drwxrwxr-x 2 seokin seokin 4096  9月 12 01:50 lab11
drwxrwxr-x 3 seokin seokin 4096  9月 12 01:50 lab12
drwxrwxr-x 3 seokin seokin 4096 10月 17 02:30 lab2
drwxrwxr-x 2 seokin seokin 4096 10月 17 04:32 lab3
drwxrwxr-x 3 seokin seokin 4096 10月 16 23:26 lab4
drwxrwxr-x 2 seokin seokin 4096 10月 24 09:27 lab5
drwxrwxr-x 2 seokin seokin 4096 10月 24 07:55 lab6
drwxrwxr-x 3 seokin seokin 4096 10月 31 00:42 lab7
drwxrwxr-x 2 seokin seokin 4096 10月 31 00:40 lab8
drwxrwxr-x 2 seokin seokin 4096  9月 12 01:50 lab9
drwxrwxr-x 2 seokin seokin 4096  9月 12 01:50 lab_dir_create
drwxrwxr-x 2 seokin seokin 4096 10月  9 19:17 mylab5
-rwxr-xr-x 1 seokin seokin 3222  9月 12 01:50 utmp.h
```

- Where is the second message?

## Example 1-1

---

- Compare pid before execvp and after
- pid\_t getpid(void)
  - returns the process ID of the current process
  - unistd.h



# Example 1-2

---

## ■ Process ID after execvp()

```
/* after.c */

#include <unistd.h>
#include <stdio.h>

void main()
{
    pid_t pid = getpid();
    printf("After execvp(): %d\n", pid);
    return;
}
```

```
/* before.c */

#include <unistd.h>
#include <stdio.h>

void main()
{
    char *arg[2];
    arg[0] = "./after";
    arg[1] = 0;

    pid_t pid = getpid();
    printf("before execvp(): %d\n", pid);
    execvp(arg[0], arg);
    return;
}
```

## Example 1-2

---

### ■ result

```
[seokin@compasslab1:~/system_programming/labs/lab7$ ./before
Before execvp(): 24317
After execvp(): 24317
[seokin@compasslab1:~/system_programming/labs/lab7$ ./before
Before execvp(): 24318
After execvp(): 24318
[seokin@compasslab1:~/system_programming/labs/lab7$ ./before
Before execvp(): 24319
After execvp(): 24319
[seokin@compasslab1:~/system_programming/labs/lab7$ ./before
Before execvp(): 24320
After execvp(): 24320
[seokin@compasslab1:~/system_programming/labs/lab7$ ./before
Before execvp(): 24321
After execvp(): 24321
```

# What does “execvp()’ system call do

---

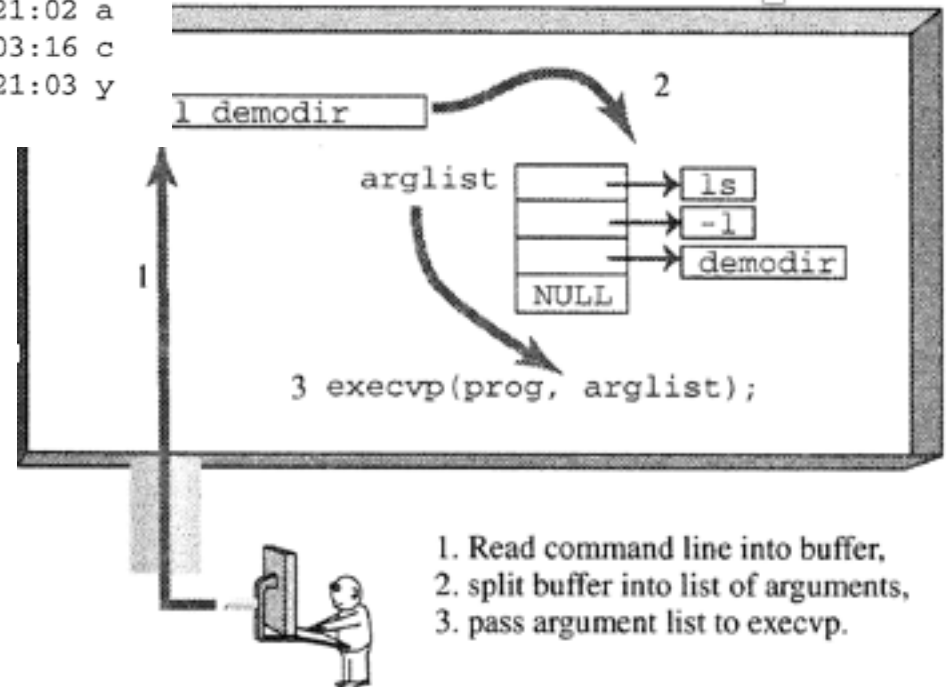
- execvp() system call clears out the machine-language code of the current program from the current process
- Puts the code of the program named in the execvp() system call
- Run that new program
- execvp() changes the memory allocation of the process to fit the space requirements of the new program.
- execvp() does not return if it succeeds
- **The process is the same, the contents are new!**

## Example 2: a prompting shell (psh1.c)

### ■ The first version of a shell

- It will prompts the user for a program name and argumetns and then run the program

```
Arg[0]? ls Enter
Arg[1]? -l Enter
Arg[2]? demodir Enter
Arg[3]? Enter
total 2
drwxr-x---  2 bruce  users 1024 Jul 14 21:02 a
drwxr-x---  3 bruce  users 1024 Jul 16 03:16 c
-rw-r--r--  1 bruce  users   0 Jul 14 21:03 y
$
```



## Example 2: a prompting shell (psh1.c)

---

```
/*      prompting shell version 1      (psh1.c)
 *
 *      Prompts for the command and its arguments.
 *      Builds the argument vector for the call to execvp.
 *      Uses execvp(), and never returns.
 */

#include      <stdio.h>
#include      <signal.h>
#include      <string.h>
#include      <stdlib.h>

#define MAXARGS      20          /* cmdline args */
#define ARGLEN      100        /* token length */

char* makestring(char*);
int execute(char*[]);
```

## Example 2: a prompting shell (psh1.c)

---

```
int main()
{
    char    *arglist[MAXARGS+1];          /* an array of ptrs    */
    int     numargs = 0;                   /* index into array    */
    char    argbuf[ARGLEN];                /* read stuff here     */

    while ( numargs < MAXARGS )
    {
        printf("Arg[%d]? ", numargs);
        if ( fgets(argbuf, ARGLEN, stdin) && *argbuf != '\n' )
            arglist[numargs++] = makestring(argbuf);
        else //입력 첫 문자가 '\n'일 때
        {
            if ( numargs > 0 ){
                arglist[numargs]=NULL;      /* any args?          */
                execute( arglist );         /* close list         */
                numargs = 0;                 /* do it              */
            }                               /* and reset          */
        }
    }
    return 0;
}
```

## Example 2: a prompting shell (psh1.c)

---

```
int execute( char *arglist[] )
/*
 *      use execvp to do it
 */
{
    execvp(arglist[0], arglist);           /* do it */
    perror("execvp failed");
    exit(1);
}
```

## Example 2: a prompting shell (psh1.c)

---

```
char * makestring( char *buf )
/*
 * trim off newline and create storage for the string
 */
{
    char    *cp ;

    buf[strlen(buf)-1] = '\0';           /* trim newline */
    cp = malloc( strlen(buf)+1 );        /* get memory    */
    if ( cp == NULL ) {                  /* or die        */
        fprintf(stderr, "no memory\n");
        exit(1);
    }
    strcpy(cp, buf);                     /* copy chars    */
    return cp;                           /* return ptr    */
}
```



## Example 2: a prompting shell (psh1.c)


---

- How'd it Go?
  - The program works OK
  - but `execvp` replaces the code of the shell with the code of the command, then exits,
- A solution is to create a new process and have that new process execute the program.


## Q2: How Do We Get a New Process?

- Ans: A process calls `fork()` system call to replicate itself.

Before fork:



After fork:



The new process contains the same code and data as the parent process.


FIGURE 8.8

`fork()` makes a copy of a process.

# fork() system call

---


- The process contains the program and a current location in the program.
- The process then calls fork().
- Control passes into the fork code inside the kernel.
- The kernel does this:
  - (a) **Allocates** a new chunk of memory and kernel data structures.
  - (b) **Copies** the original process into the new process
  - (c) **Adds** the new process to the set of running processes
  - (d) **Returns** control back to both processes



# fork() system call

---

- After a process calls fork, there are two separate process,
  - both digitally identical,
  - both in the middle of the same code,
  - but each a separate process can go its own way.



# Creating an new process Example : forkdemo1.c

---

```
/*  forkdemo1.c
 *      shows how fork creates two processes, distinguishable
 *      by the different return values from fork()
 */
#include      <stdio.h>
#include      <unistd.h>
main()
{
    int      ret_from_fork, mypid;

    mypid = getpid();          /* who am i?          */
    printf("Before: my pid is %d\n", mypid);    /* tell the world */

    ret_from_fork = fork();

    sleep(1);
    printf("After: my pid is %d, fork() said %d\n",
           getpid(), ret_from_fork);
}
```

```
$ cc forkdemo1.c -o forkdemo1
```

```
$ ./forkdemo1
```

```
Before: my pid is 4170
```


```
After: my pid is 4170, fork() said 4171
```

```
$ After: my pid is 4171, fork() said 0
```

Three lines!... Why?


# Creating an new process Example : forkdemo1.c

Before fork:



One flow of control enters the fork kernel code.

After fork:



Two flows of control return from fork kernel code.

FIGURE 8.9

The child executes the code after `fork()`.

# Children creating process example : forkdemo2.c


---

```
/* forkdemo2.c - shows how child processes pick up at the return
 *                from fork() and can execute any code they like,
 *                even fork().  Predict number of lines of output.
 */
#include          <stdio.h>
#include          <unistd.h>
main()
{
    printf("my pid is %d\n", getpid() );
    fork();
    fork();
    fork();
    printf("my pid is %d\n", getpid() );
}
```

# Children creating process example : forkdemo2.c

---

```
my pid is 24740
my pid is 24740
my pid is 24741
my pid is 24742
my pid is 24744
my pid is 24745
my pid is 24743
my pid is 24746
my pid is 24747
seokin@compasslab1:~/system_programming/labs/lab7$
```





# fork() system call

---

fork	
PURPOSE	Create a process
INCLUDE	#include < unistd.h >
USAGE	pid_t result = fork(void)
ARGS	none
RETURNS	<div>-1 : if error 0 : to child process pid : pid of child to parent process</div>

# Distinguishing parent from child Example : forkdemo3.c

---

```
/* forkdemo3.c - shows how the return value from fork()
 *               allows a process to determine whether
 *               it is a child or process
 */
#include        <stdio.h>
#include        <unistd.h>
main()
{
    int         fork_rv;
    printf("Before: my pid is %d\n", getpid());

    fork_rv = fork();                    /* create new process */

    if ( fork_rv == -1 )                  /* check for error */
        perror("fork");

    else if ( fork_rv == 0 )
        printf("I am the child.  my pid=%d\n", getpid());
    else
        printf("I am the parent. my child is %d\n", fork_rv);
}
```

---

■ Three skills needed to build a shell:

- 1. Create a new process
- 2. Run a program (execvp)
- **3. Tell the parent wait until the child process finishes executing the command**

### Q3: How Does the Parent Wait for the Child to exit?

---

- Ans: A process calls wait() system call to wait for a child to finish.

pid = wait( &status );

- The wait system call does two things:
  - It pauses the calling process until a child process finishes running.
  - It retrieves the value the child process had passed to exit.

# wait() system call

---

<b>wait</b>	
PURPOSE	Wait for process termination
INCLUDE	<code>#include &lt; sys/ types.h &gt;</code> <code>#include &lt;sys/ wait.h &gt;</code>
USAGE	<code>pid_t result = wait( int *statusptr)</code>
ARGS	<code>statusptr</code> <code>child</code> <code>result</code>
RETURNS	<code>-1</code> : if error <code>pid</code> : of terminated process
SEE ALSO	<code>waitpid(2)</code> , <code>wait3(2)</code>

# wait() system call

---

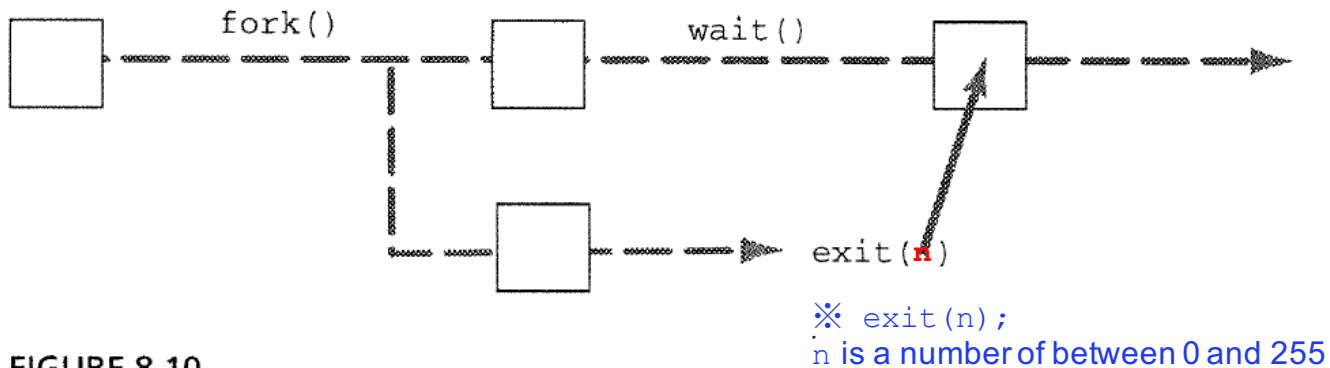


FIGURE 8.10

`wait` pauses the parent until the child finishes.

- When the child calls `exit()`,
  - the kernel wakes up the parent and
  - delivers the value the child passed to `exit()`.
- Thus, `wait()` performs two operations
  - notification and communication

# Notification Example : waitdemo1.c

---

```
/* waitdemo1.c - shows how parent pauses until child finishes
*/

#include <stdio.h>
#include <stdlib.h>
#define DELAY 2

main()
{
    int newpid;
    void child_code(int), parent_code(int);
    printf("before: mypid is %d\n", getpid());

    if ( (newpid = fork()) == -1 )
        perror("fork");
    else if ( newpid == 0 )
        child_code(DELAY);
    else
        parent_code(newpid);
}
```

## Notification Example : waitdemo1.c

---

```
/*  
 * new process takes a nap and then exits  
 */  
void child_code(int delay)  
{  
    printf("child %d here. will sleep for %d seconds\n", getpid(),  
        delay);  
    sleep(delay);  
    printf("child done. about to exit\n");  
    exit(17);  
}
```



# Notification Example : waitdemo1.c

---

```
/*
 * parent waits for child then prints a message
 */
void parent_code(int childpid)
{
    int wait_rv;          /* return value from wait() */
    wait_rv = wait(NULL);
    printf("done waiting for %d. Wait returned: %d\n", childpid,
        wait_rv);
}
```

```
$ ./waitdemo1
before: mypid is 10328
child 10329 here. will sleep for 2 seconds
child done. about to exit
done waiting for 10329. Wait returned: 10329
```

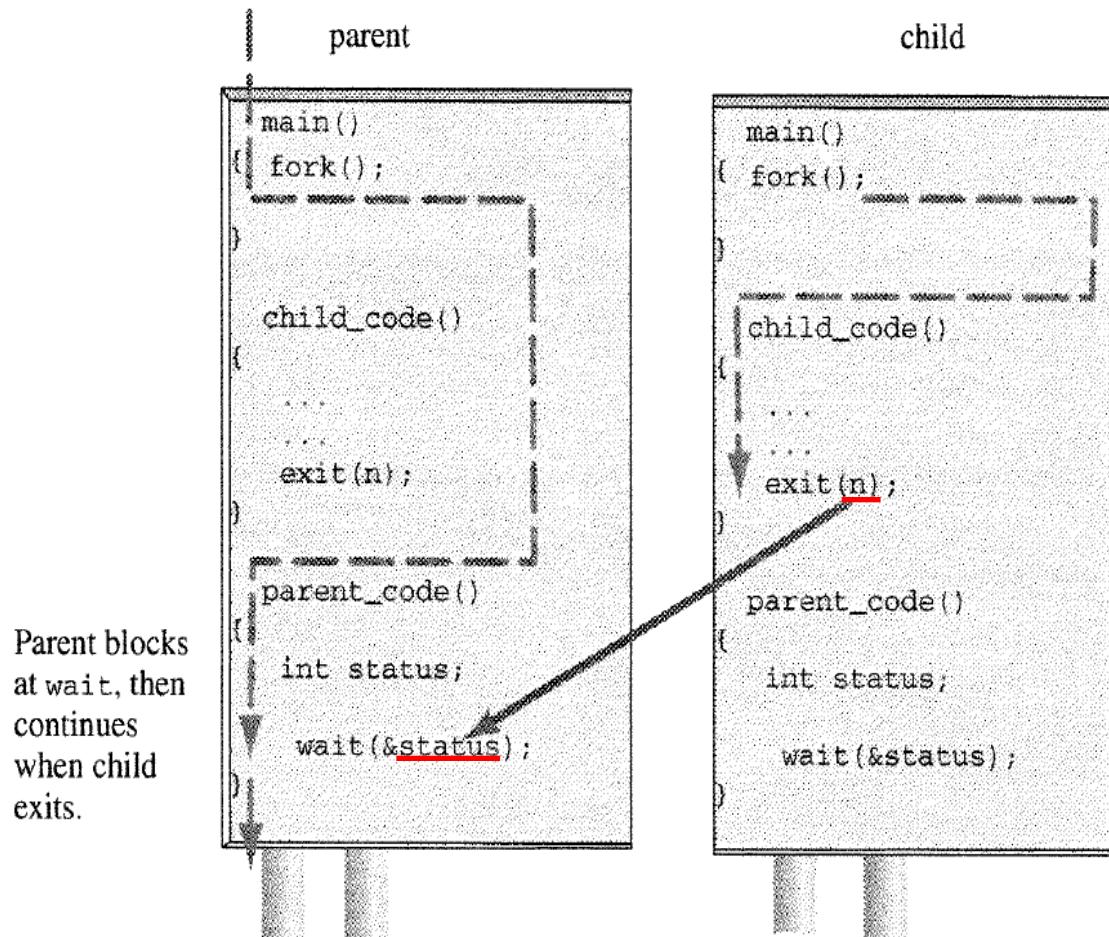


FIGURE 8.11

Control flow and communication with `wait()`.

# Purpose of wait() system call

---

- To notify the parent that a child process finished running
- To tell the parent how a child process finished
  - A process ends in one of three ways
    - Success
    - Failure
    - Death
  - A process can succeed at its task:  
    `exit(0)` or `return 0` from `main`
  - A process can fail at its task:  
    `exit(nonzerovalue)`
  - A process might be killed by a signal

# How does the parent know if the child's exit status

- How does the parent know if the child succeeded, failed, or died?
  - Answer is in the argument (integer pointer) to wait() system call
    - Eight bit for exit value, seven bits for signal number, one bit to indicate a core dump




FIGURE 8.12

The child status value has three parts.

# Communication Example: waitdemo2.c

---

```
/* waitdemo2.c - shows how parent gets child status
*/

#include <stdio.h>
#include <stdlib.h>
#define DELAY 5

main()
{
    int newpid;
    void child_code(), parent_code();
    printf("before: mypid is %d\n", getpid());
    if ( (newpid = fork()) == -1 )
        perror("fork");
    else if ( newpid == 0 )
        child_code(DELAY);
    else
        parent_code(newpid);
}
```

## Communication Example: waitdemo2.c

---

```
/*
 * new process takes a nap and then exits
 */
void child_code(int delay)
{
    printf("child %d here. will sleep for %d seconds\n", getpid(),
        delay);
    sleep(delay);
    printf("child done. about to exit\n");
    exit(17);
}
```

# Communication Example: waitdemo2.c

---

```
/*
 * parent waits for child then prints a message
 */
void parent_code(int childpid)
{
    int wait_rv;           /* return value from wait() */
    int child_status;
    int high_8, low_7, bit_7;

    wait_rv = wait(&child_status);
    printf("done waiting for %d. Wait returned: %d\n", childpid,
        wait_rv);

    high_8 = child_status >> 8;    /* 1111 1111 0000 0000 */
    low_7  = child_status & 0x7F;   /* 0000 0000 0111 1111 */
    bit_7  = child_status & 0x80;   /* 0000 0000 1000 0000 */
    printf("status: exit=%d, sig=%d, core=%d\n", high_8, low_7,
        bit_7);
}
```

# Communication Example: waitdemo2.c

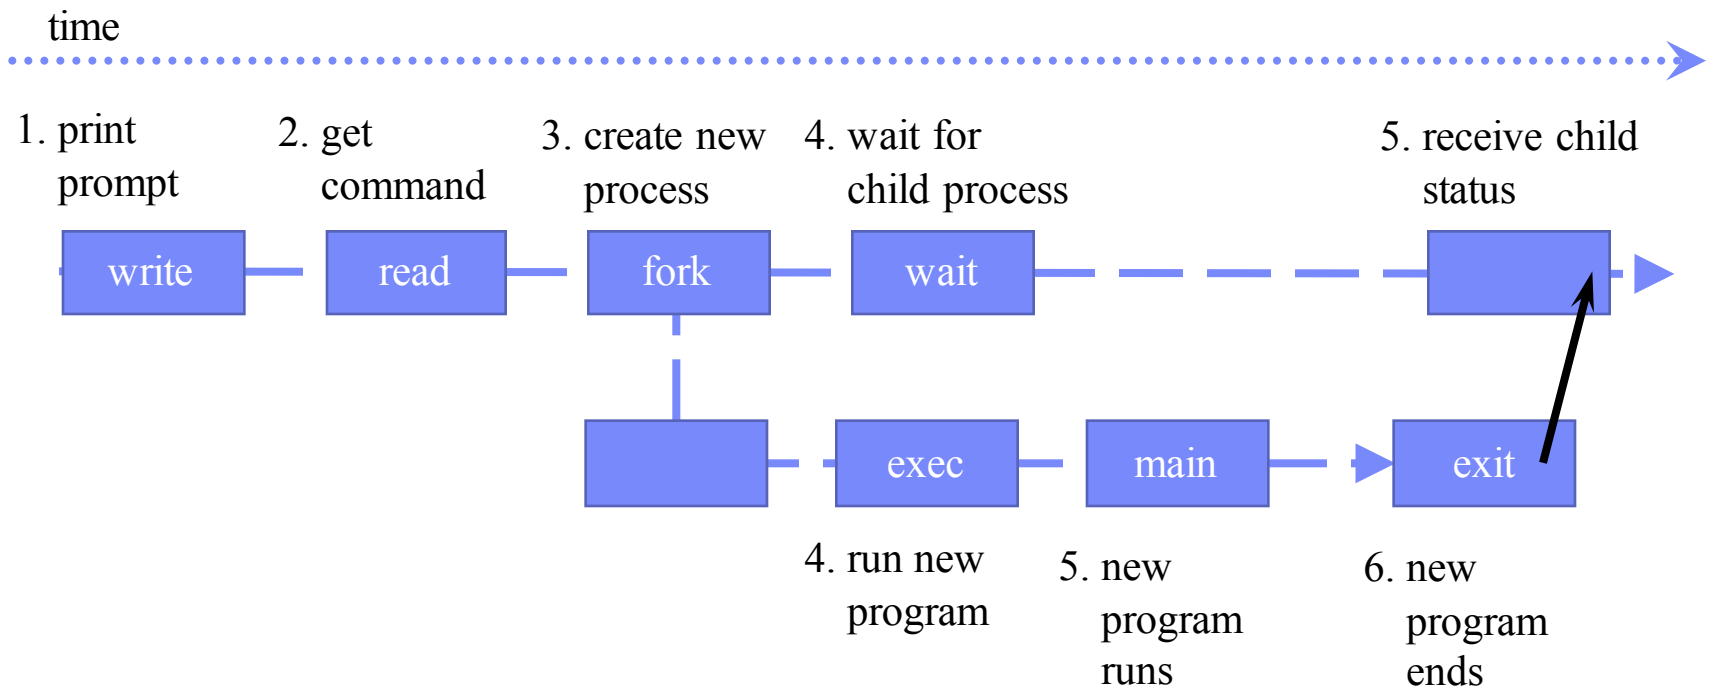
---

- Run in the background and use kill to send SIGTERM to the child:

```
$ ./waitdemo2 &  
$ before: mypid is 10857  
child 10858 here. will sleep for 5 seconds  
kill 10858  
$ done waiting for 10858. Wait returned: 10858  
status: exit=0, sig=15, core=0  
SIGTERM
```



# 8.4.5 Summary : How the Shell Runs Programs



**FIGURE 8.13** Shell loop with fork(), exec(), and wait()

# Contents

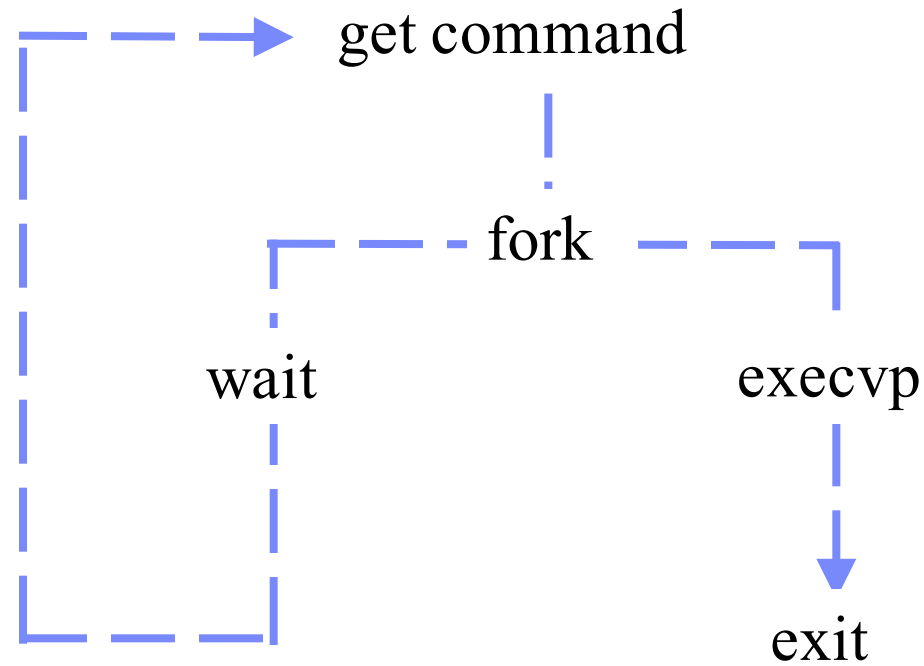
---

- 8.1 Process
- 8.2 Learning about Processes with `ps`
- 8.3 The Shell: A Tool for Process and Program Control
- 8.4 How the Shell Runs Programs
- 8.5 Writing a Shell: `psh2.c`

# Wring a shell : psh2.c

---

- Logic diagram



# Wring a shell : psh2.c

---

```
/** prompting shell version 2 (psh2.c)
**
**          Solves the 'one-shot' problem of version 1
**          Uses execvp(), but fork()s first so that the
**          shell waits around to perform another command
**          New problem: shell catches signals.  Run vi, press ^c.
**/

#include      <stdio.h>
#include      <signal.h>
#include      <string.h>
#include      <stdlib.h>

#define MAXARGS      20          /* cmdline args */
#define ARGLEN       100        /* token length */

char* makestring(char*);
void execute(char*[]);
```

# Wring a shell : psh2.c

---

```
int main()
{
    char    *arglist[MAXARGS+1];          /* an array of ptrs    */
    int     numargs = 0;                   /* index into array    */
    char     argbuf[ARGLEN];               /* read stuff here     */

    while ( numargs < MAXARGS )
    {
        printf("Arg[%d]? ", numargs);
        if ( fgets(argbuf, ARGLEN, stdin) && *argbuf != '\n' )
            arglist[numargs++] = makestring(argbuf);
        else
        {
            if ( numargs > 0 ){              /* any args?          */
                arglist[numargs]=NULL;      /* close list         */
                execute( arglist );         /* do it              */
                numargs = 0;                 /* and reset          */
            }
        }
    }
    return 0;
}
```

## Wring a shell : psh2.c

```
-execute( char *arglist[] )  
/*  
 *    use fork and execvp and wait to do it  
 */  
{  
    int    pid,exitstatus;                /* of child    */  
    pid = fork();                        /* make new process */  
    switch( pid ){  
        case -1:  
            perror("fork failed");  
            exit(1);  
        case 0:  
            execvp(arglist[0], arglist);    /* do it */  
            perror("execvp failed");  
            exit(1);  
        default:  
            while( wait(&exitstatus) != pid )  
                ;  
            printf("child exited with status %d,%d\n",  
                    exitstatus>>8, exitstatus&0377);  
    }  
}
```

## Wring a shell : psh2.c

---

```
char *makestring( char *buf )
/*
 * trim off newline and create storage for the string
 */
{
    char      *cp, *malloc();

    buf[strlen(buf)-1] = '\0';           /* trim newline */
    cp = malloc( strlen(buf)+1 );       /* get memory */
    if ( cp == NULL ){                  /* or die */
        fprintf(stderr, "no memory\n");
        exit(1);
    }
    strcpy(cp, buf);                    /* copy chars */
    return cp;                          /* return ptr */
}
```

---

\$ ./psh2

Arg[0]? **tr**

Arg[1]? **[a-z]**

Arg[2]? **[A-Z]**

Arg[3]?

**hello**

HELLO

**now to press**

NOW TO PRESS

**Ctrl-C***press ^C here*

\$




# Signals and psh2.c

---

- What happen if we press Ctrl-C when psh2 is waiting for the child process to finish? ...

```
$ ./psh2
Arg[0]? tr
Arg[1]? [a-z]
Arg[2]? [A-Z]
Arg[3]?
hello
HELLO
now to press
NOW TO PRESS
Ctrl-Cpress ^C here
$
```



The SIGINT killed not only 'tr' but also 'psh2'