

Ch8. I/O Redirection and Pipes

Prof. Seokin Hong

Kyungpook National University

Fall 2018

THE SHELL

- A shell is a program that manages processes and runs programs
- Thee main functions of shells
 - (a) Shells run programs
 - (b) Shells manage input and output
 - (c) Shells can be programmed

THE SHELL

- How do the following commands work?

```
$ ls > my.files
```

```
$ who | sort > userlist
```

- We focus on a particular form of interprocess communication: **I/O redirection and pipes**

Objectives

■ Ideas and Skills

- I/O Redirection : What and why?
- Definitions of standard input, output, and error
- Redirecting standard I/O to files
- Using fork to redirect I/O for other programs
- **Pipes**
- Using fork with pipes

■ System Calls and Functions

- dup, dup2
- pipe

■ Consider the following problem:

You want a program that notifies you when people log in or log out of the system so you can watch for your pals.

- You could write a C program that uses the utmp file and interval timers.
- A simpler solution is to write a shell script : who

A SHELL APPLICATION : WATCH FOR USERS

Logic

Get list of users (call it prev)

While true

 sleep

 get list of users (call it curr)

 compare lists

 in prev, not in curr -> logout

 in curr, not in prev -> login

 make prev = curr

repeat

Shell code

Who | sort > prev

While true ; do

 sleep 60

 who | sort > curr

 echo "logged out:"

 comm -23 prev curr

 echo "logged in:"

 comm -13 prev curr

 mv curr prev

done


A SHELL APPLICATION : WATCH FOR USERS

shell code

■ `who | sort > prev`

```
who | sort > prev
while true ; do
    sleep 60
    who | sort > curr
    echo "logged out:"
    comm -23 prev curr
    echo "logged in:"
    comm -13 prev curr
    mv curr prev
done
```

`who` | `sort` > `file`



A SHELL APPLICATION : WATCH FOR USERS

■ The comm command

- compares two sorted lists
- prints out three columns: 1, 2, and 3.

```
root@DESKTOP-K4MA2V5:~# cat file1.txt
```

```
a  
b  
c  
d
```

```
root@DESKTOP-K4MA2V5:~# cat file2.txt
```

```
a  
b  
c  
e
```


```
root@DESKTOP-K4MA2V5:~# comm file1.txt file2.txt
```

```
d  
e  
a  
b  
c
```

```
root@DESKTOP-K4MA2V5:~#
```

shell code

```
-----  
who | sort > prev  
while true ; do  
    sleep 60  
    who | sort > curr  
    echo "logged out:"  
    comm -23 prev curr  
    echo "logged in:"  
    comm -13 prev curr  
    mv curr prev  
done
```



A SHELL APPLICATION : WATCH FOR USERS

- `comm -23 prev curr`
 - drop columns 2 and 3 → show lines only in prev
- `comm -13 prev curr`
 - drop columns 1 and 3 → show lines only in curr




FIGURE 10.2

`comm` compares two lists and outputs three sets.

```
shell code
-----
who | sort > prev
while true ; do
    sleep 60
    who | sort > curr
    echo "logged out:"
    comm -23 prev curr
    echo "logged in:"
    comm -13 prev curr
    mv curr prev
done
```

Contents

- 10.1 Shell
- 10.2 A Shell Application : Watch for Users
- 10.3 Facts about Standard I/O and Redirection
- 10.4 How to Attach stdin to a File
- 10.5 Redirecting I/O for Another Program: `who>userlist`
- 10.6 Programming Pipes

- Lessons:

- (a) Power of shell scripts—easier and quicker than C
- (b) Flexibility of software tools—each tool does one specific, general task
- (c) Use and value of I/O redirection and pipes

- The **>** operator can be used to treat *files as variables* of arbitrary size and structure.


`x = func_a(func_b(y));` in C

`prog_b | prog_a > x` in sh

- Question: How?


FACT ABOUT STANDARD I/O AND REDIRECTION

- All unix I/O redirection is based on the principle of **standard streams of data**.
- **Standard input**
 - the stream of data to process
- **Standard output**
 - the stream of result data
- **Standard error**
 - a stream of error messages



Fact One: Three Standard File Descriptors

- All Unix tools use file descriptor 0, 1, and 2.
 - 0: standard in (stdin)
 - 1: standard out (stdout)
 - 2: standard error (stderr)
- Every unix program gets three open file descriptors at startup:



Output Goes Only to stdout

- Most programs do NOT accept names for output files;
 - They always write **results** to **file descriptor 1** and **errors** to **file descriptor 2**.
 - If you want to send the output of a process to a file or to the input of another process, you need to change where the file descriptor goes.

The Shell, Not the Program, Redirects I/O

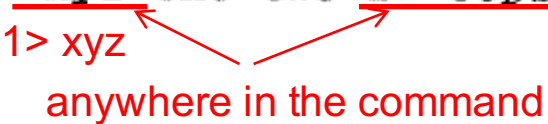
- You **tell the shell to attach file descriptor 1 to a file** by using the output redirection notation:
 \$ cmd > filename
- The shell **connects that file descriptor to the named file.**
- The **program continues to write to file descriptor 1, unaware of the new data destination.**

```
/* listargs.c
 *          print the number of command line args, list the args,
 *          then print a message to stderr
 */
#include    <stdio.h>

main( int ac, char *av[] )
{
    int     i;

    printf("Number of args: %d, Args are:\n", ac);
    for(i=0;i<ac;i++)
        printf("args[%d] %s\n", i, av[i]);           // to stdout
    fprintf(stderr,"This message is sent to stderr.\n"); // to stderr
}
```

```
$ cc listargs.c -o listargs
$ ./listargs testing one two
args[0] ./listargs
args[1] testing
args[2] one
args[3] two
This message is sent to stderr.
$ ./listargs testing one two > xyz
This message is sent to stderr.
$ cat xyz
args[0] ./listargs
args[1] testing
args[2] one
args[3] two
$ ./listargs testing >xyz one two 2> oops
$ cat xyz
args[0] ./listargs
args[1] testing
args[2] one
args[3] two
$ cat oops
This message is sent to stderr.
```



1> xyz
anywhere in the command

Understanding I/O Redirection

■ Goal:

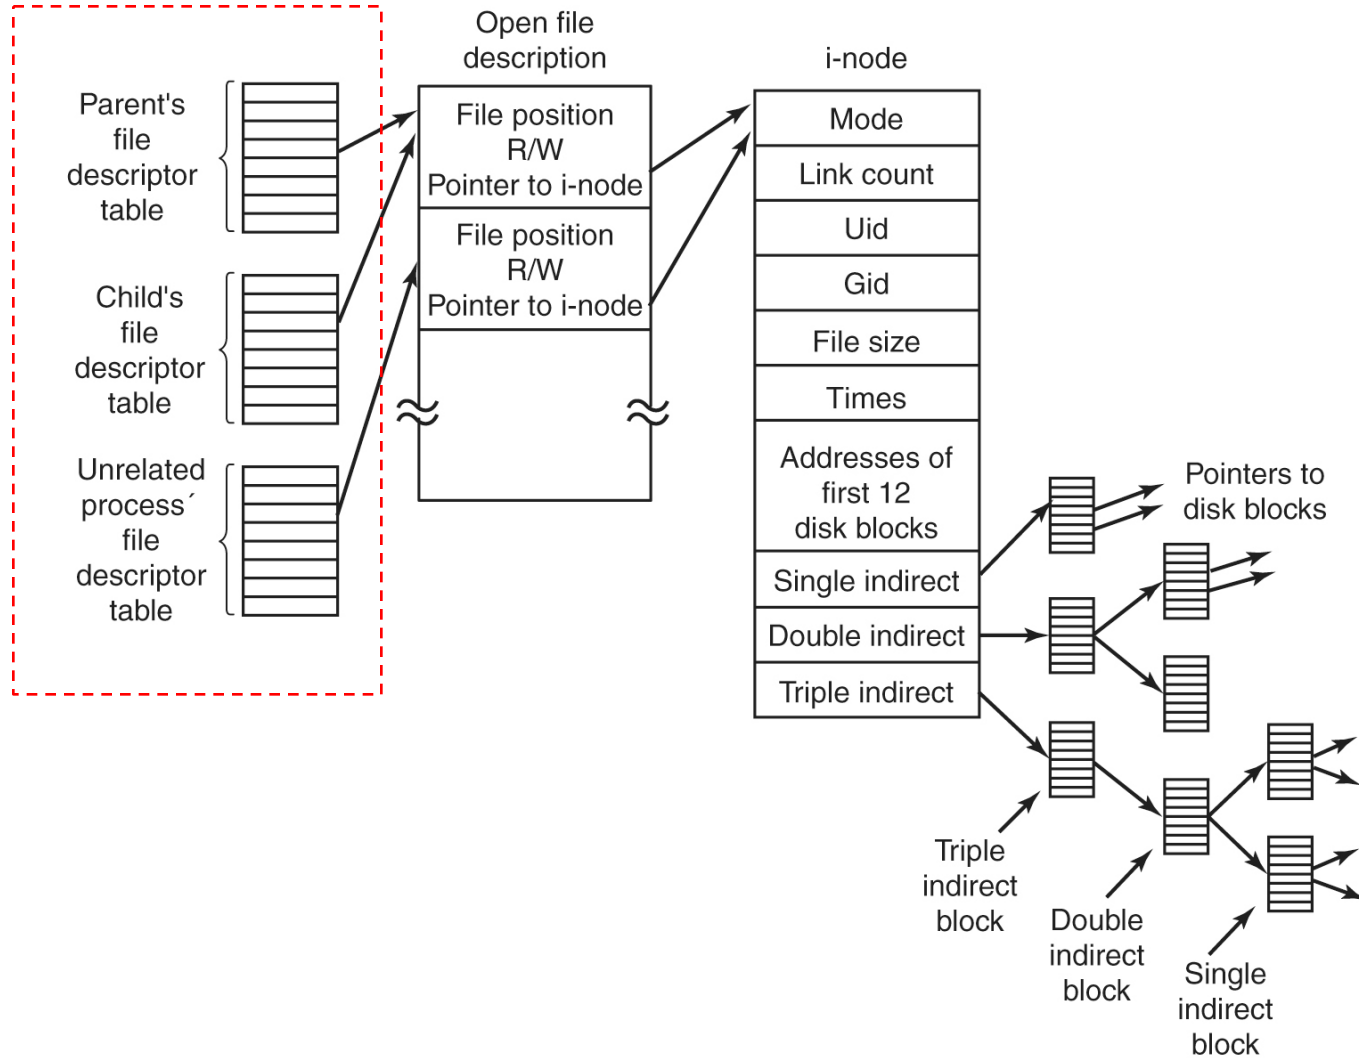
- Understand how I/O redirection works
- Learn how to write programs that use it

who > userlist	attach stdout to a file
sort < data	attach stdin to a file
who sort	attach stdout to stdin

Fact Two: The “Lowest-Available-fd” Principle

- What is a **file descriptor**? It is an index of an array.
 - Each process has a collection of files it has open.
 - The information of those open files are kept in an array.
- **FACT:** When you open a file, you always get the lowest available spot in the array.

File descriptor



Contents

- 10.1 Shell
- 10.2 A Shell Application : Watch for Users
- 10.3 Facts about Standard I/O and Redirection
- 10.4 How to Attach stdin to a File
- 10.5 Redirecting I/O for Another Program: `who>userlist`
- 10.6 Programming Pipes

How to Attach stdin to a File

- How does a program redirect standard input so that data come from a file?

ex) `$ sort < data`

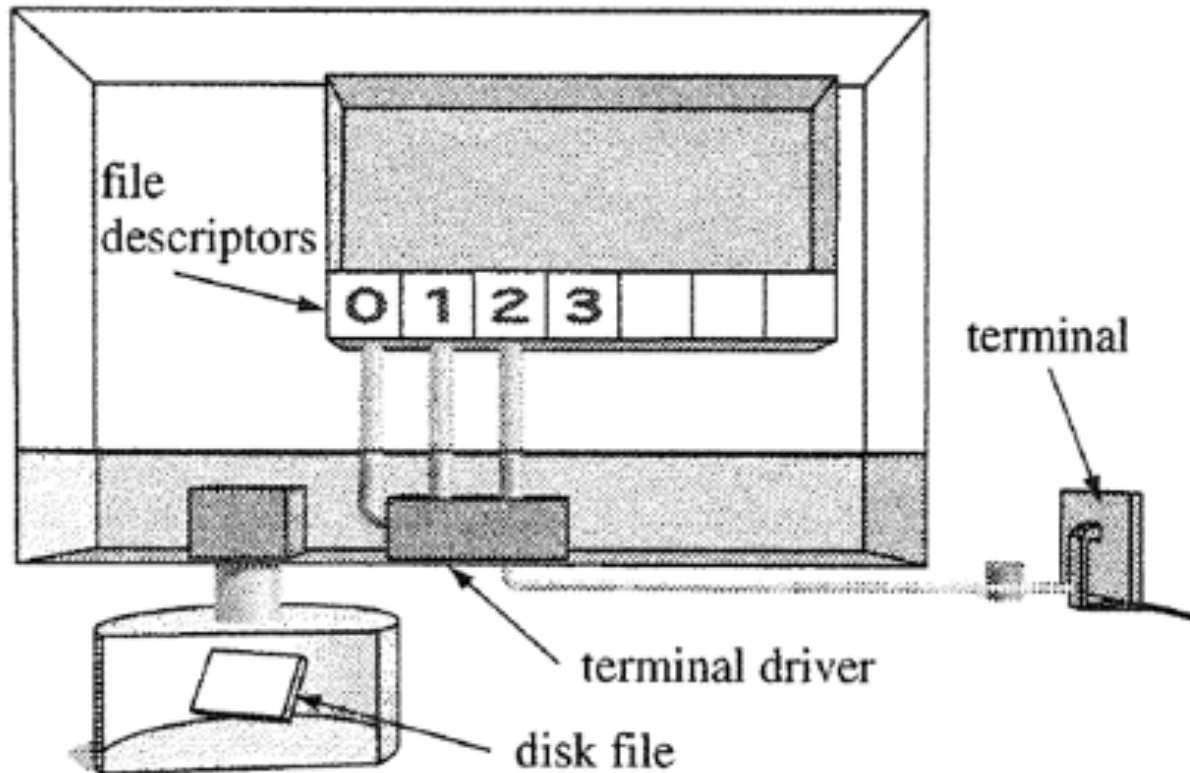
```
root@DESKTOP-K4MA2V5:~# cat data
5
4
3
2
1
root@DESKTOP-K4MA2V5:~# sort < data
1
2
3
4
5
root@DESKTOP-K4MA2V5:~#
```

- If we attach file descriptor 0 to a file, that file becomes the source for standard input.
- How? ...

Method 1: Close Then Open

■ Starting


- File descriptor 0, 1, 2 attached to the terminal driver



Method 1: Close Then Open

- **Then, close(0)**


- The first element in the array of file descriptors is now unused



Method 1: Close Then Open

■ Finally, `open(filename, O_RDONLY)`

- Opens the file you want to attach to stdin.



```

/* stdinredir1.c      close-then-open method
 *
 *      purpose: show how to redirect standard input by replacing file
 *
 *      descriptor 0 with a connection to a file.
 *
 *      action: reads three lines from standard input, then
 *
 *      closes fd 0, opens a disk file, then reads in
 *
 *      three more lines from standard input
#include      <unistd.h>
#include      <stdio.h>
#include      <fcntl.h>
#include      <stdlib.h>
main()
{

    int      fd ;
    char      line[100];

    /* read and print three lines */

    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );

    /* redirect input */

    close(0);
    fd = open("/etc/passwd", O_RDONLY);
    if ( fd != 0 ){
        fprintf(stderr, "Could not open data as fd 0\n");
        exit(1);
    }

    /* read and print three lines */

    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
}

```


Method 2: open..close..**dup**..close

Method 3: open..**dup2**..close


dup, dup2	
PURPOSE	Copy a file descriptor
INCLUDE	#include <unistd.h>
USAGE	newfd = dup(oldfd); newfd = dup2(oldfd, newfd);
ARGS	oldfd file descriptor to copy newfd copy of oldfd
RETURNS	-1 if error newfd new file descriptor

```
#ifdef CLOSE_DUP
    close(0);
    newfd = dup(fd);
#else
    newfd = dup2(fd, 0);
#endif
```


```
fd = open("f", O_RDONLY);
```



```
close(0);
```



```
dup(fd);
```



```
close(fd);
```




FIGURE 10.9

Using `dup` to redirect.

-
- open(file)* The first step is to open the file to which `stdin` should be attached. This call returns a file descriptor, which is not 0, since 0 is currently open.
- close(0)* The next step is to close file descriptor 0. File descriptor 0 is now unused.
- dup(fd)* The `dup(fd)` system call makes a duplicate of `fd`. The duplicate uses the lowest number unused file descriptor. Therefore, the duplicate of the connection to the file is at spot 0 in the array of open files. We have thereby attached the disk file to file descriptor 0.
- close(fd)* Finally, we `close(fd)`, the original connection to the file, leaving only the connection on file descriptor 0. Compare this method to the technique of moving a phone call from one extension to another.

```
/* stdinredir2.c
 *      shows two more methods for redirecting standard input
 *      use #define to set one or the other
 */
#include <unistd.h>
#include <stdio.h>
#include <fcntl.h>
#include <stdlib.h>
/* #define CLOSE_DUP      /* open, close, dup, close */
/* #define USE_DUP2      /* open, dup2, close */

main()
{
    int    fd ;
    int    newfd;
    char    line[100];

    /* read and print three lines */

    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
```

```
    /* redirect input */ ※ test with "/etc/passwd"
    fd = open("data", O_RDONLY);    /* open the disk file    */
#ifdef CLOSE_DUP
    close(0);
    newfd = dup(fd);                /* copy open fd to 0    */
#else
    newfd = dup2(fd, 0);            /* close 0, dup fd to 0 */
#endif
    if ( newfd != 0 ){
        fprintf(stderr, "Could not duplicate fd to 0\n");
        exit(1);
    }
    close(fd);                      /* close original fd    */

    /* read and print three lines */

    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
}
```

But the Shell Redirects stdin for Other Programs

- In practice, of course, if a program wants to read a file, it can just open the file directly rather than changing standard input
- The real value of these samples is to show how one program can change standard input for another program;
\$ sort < data

Contents

- 10.1 Shell Programming
- 10.2 A Shell Application : Watch for Users
- 10.3 Facts about Standard I/O and Redirection
- 10.4 How to Attach stdin to a File
- 10.5 Redirecting I/O for Another Program: `who > userlist`
- 10.6 Programming Pipes

Redirecting I/O for other programs

- The shell redirects output for a child.
 - How “who > userlist” works?

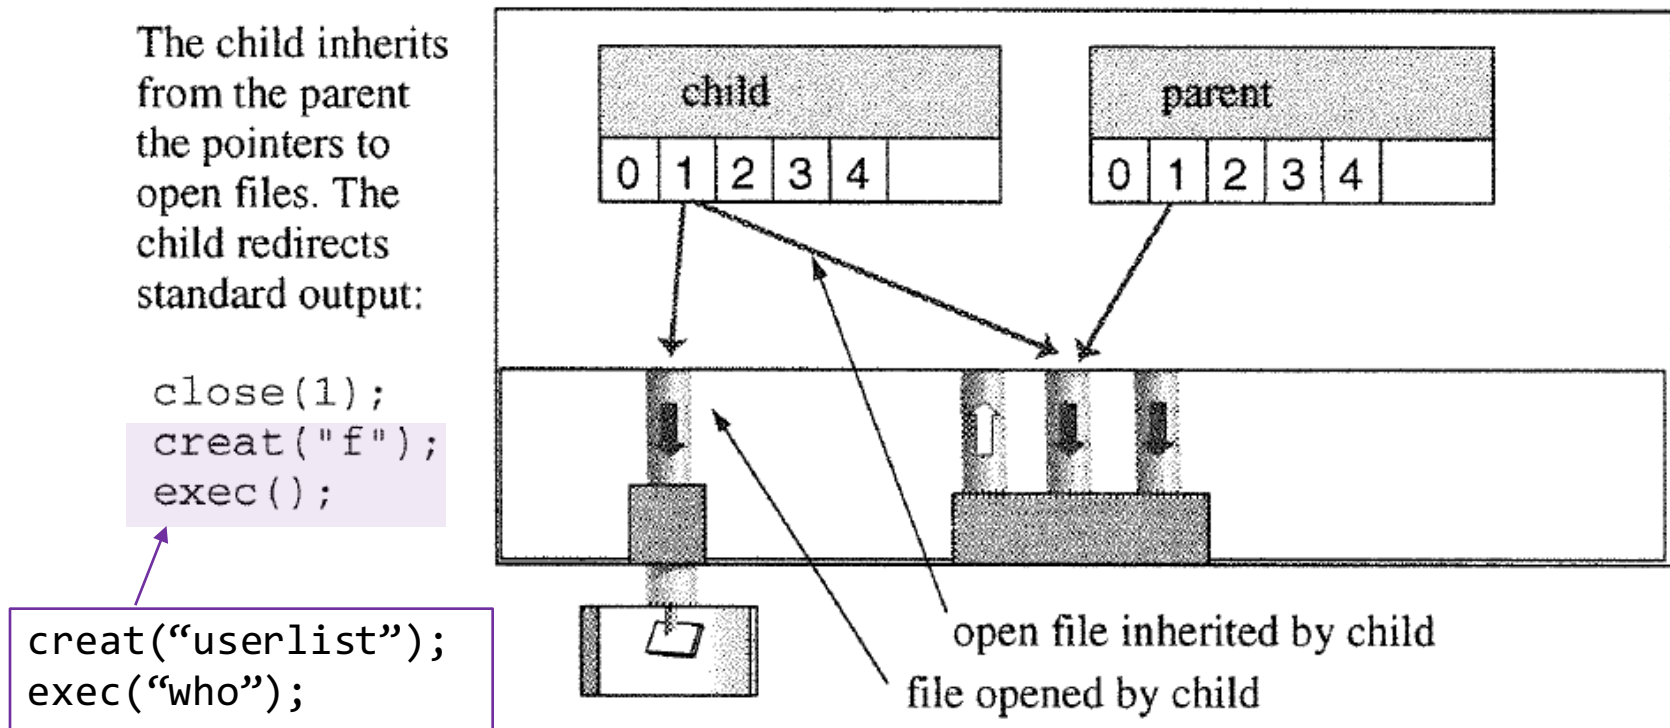
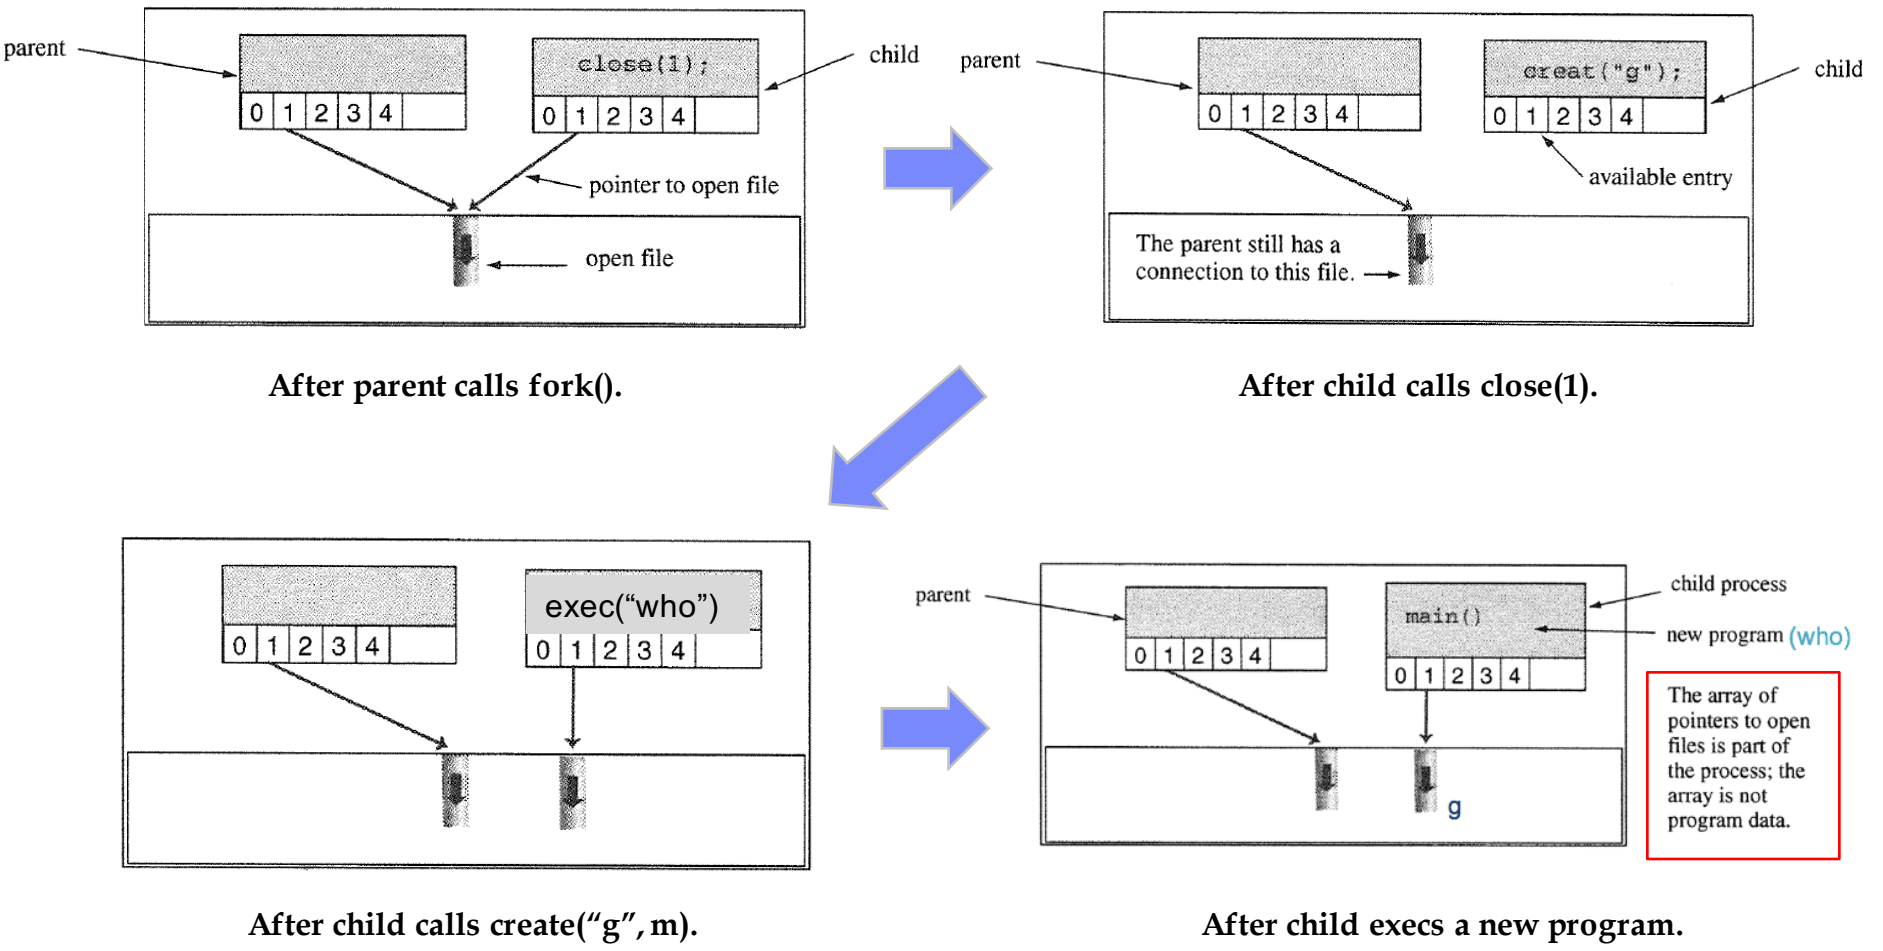


FIGURE 10.10

The shell redirects output for a child.

Redirecting I/O for other programs

■ who > g



```
/* whotofile.c
 *      purpose: show how to redirect output for another program
 *      idea: fork, then in the child, redirect output, then exec
 */

#include      <stdio.h>
#include      <unistd.h>      /* for execlp */
#include      <stdlib.h>      /* for exit */
main()
{
    int      pid ;
    int      fd;

    printf("About to run who into a file\n");

    /* create a new process or quit */
    if( (pid = fork() ) == -1 ){
        perror("fork"); exit(1);
    }
    /* child does the work */
    if ( pid == 0 ){
        close(1);
        fd = creat( "userlist", 0644 );
        execlp( "who", "who", NULL );
        perror("execlp");
        exit(1);
    }
    /* parent waits then reports */
    if ( pid != 0 ){
        wait(NULL);
        printf("Done running who.  results in userlist\n");
    }
}
```

Summary of Redirection to Files

- Three basic facts:

- (a) Standard input, output, and error are file descriptors 0, 1, and 2.
- (b) The kernel always uses the lowest numbered unused file descriptor.
- (c) The set of file descriptors is passed unchanged across `exec` calls.

- The shell also supports the following forms:

- `who > userlog` → write
- `who >> userlog` → add
- `sort < data` → read

Contents

- 10.1 Shell Programming
- 10.2 A Shell Application : Watch for Users
- 10.3 Facts about Standard I/O and Redirection
- 10.4 How to Attach stdin to a File
- 10.5 Redirecting I/O for Another Program: `who > userlist`
- 10.6 Programming Pipes

Pipe



Pipe

- `$ who | sort`
- A pipe is a one-way data channel in the kernel
 - It has a reading end and a writing end.





FIGURE 10.16
Two processes connected by a pipe.

- How to create a pipe and how to connect stdin and stdout to a pipe?

Creating a Pipe

- Use `pipe()` system call:
 - It creates the pipe and connects its two ends to two file descriptors.

pipe	
PURPOSE	Create a pipe
INCLUDE	#include <unistd.h>
USAGE	result = pipe(int array[2])
ARGS	<u>array</u> <u>an array of two ints</u>
RETURNS	-1 if error 0 if success



- array[0] is the file descriptor of the reading end
- array[1] is the file descriptor of the writing end


```

int    len, i, apipe[2];      /* two
char    buf[BUFSIZ];          /* for

/* get a pipe */
if ( pipe ( apipe ) == -1 ){
    perror("could not make pipe");
    exit(1);
}


```

Before pipe



The process has some usual files open.

After pipe



The kernel creates a pipe and sets file descriptors.

FIGURE 10.18

A process creates a pipe.

※ pipe uses the lowest-numbered available file descriptors.

```
/* pipedemo.c  * Demonstrates: how to create and use a pipe
 *
 *            * Effect: creates a pipe, writes into writing
 *              end, then runs around and reads from reading
 *              end.  A little weird, but demonstrates the idea.
 */
#include      <stdio.h>
#include      <unistd.h>

main()
{
    int      len, i, apipe[2];          /* two file descriptors */
    char      buf[BUFSIZ];              /* for reading end      */

    /* get a pipe */
    if ( pipe ( apipe ) == -1 ){
        perror("could not make pipe");
        exit(1);
    }
    printf("Got a pipe! It is file descriptors: { %d %d }\n",
           apipe[0], apipe[1]);

    /* read from stdin, write into pipe, read from pipe, print */
```


```
while ( fgets(buf, BUFSIZ, stdin) ){
    len = strlen( buf );
    if ( write( apipe[1], buf, len) != len ){
        perror("writing to pipe");
        break;
    }
    for ( i = 0 ; i<len ; i++ )
        buf[i] = 'X' ;
    len = read( apipe[0], buf, BUFSIZ ) ;
    if ( len == -1 ){
        perror("reading from pipe");
        break;
    }
    if ( write( 1, buf, len ) != len ){
        perror("writing to stdout");
        break;
    }
}
}
```

/* send */
/* down */
/* pipe */

/* wipe */

/* read */
/* from */
/* pipe */

/* send */
/* to */
/* stdout */




```
while ( fgets(buf, BUFSIZ, stdin) ){
    len = strlen( buf );
    if ( write( apipe[1], buf, len) != len ){
        perror("writing to pipe");
        break;
    }
    for ( i = 0 ; i<len ; i++ )
        buf[i] = 'X' ;
    len = read( apipe[0], buf, BUFSIZ ) ;
    if ( len == -1 ){
        perror("reading from pipe");
        break;
    }
    if ( write( 1, buf, len ) != len ){
        perror("writing to stdout");
        break;
    }
}
```

FIGURE 10.19
Data flow in `pipdemo.c`.

Using fork to Share a Pipe

Sharing a pipe:

A process calls pipe. The kernel creates a pipe and adds to the array of file descriptors pointers to the ends of the pipe.



The process then calls fork. The kernel creates a new process, and copies into that process the array of file descriptors from the parent.

Both processes have access to both ends of one pipe.




FIGURE 10.20

Sharing a pipe.

- Both processes can read and write, but a pipe is most effective when one process writes data and the other process reads data;
 - If child writes into the pipe, the parent can read those bytes




FIGURE 10.21
Interprocess data flow.

```
/* pipedemo2.c  * Demonstrates how pipe is duplicated in fork()
 *              * Parent continues to write and read pipe,
 *              * but child also writes to the pipe
 */
#include        <stdio.h>

#define CHILD_MESS    "I want a cookie\n"
#define PAR_MESS      "testing..\n"
#define oops(m,x)     { perror(m); exit(x); }
```



```
main()
```

```
{
```

```
    int pipefd[2];          /* the pipe */
    int len;                /* for write */
    char buf[BUFSIZ];       /* for read */
    int read_len;

    if ( pipe( pipefd ) == -1 )
        oops("cannot get a pipe", 1);

    switch( fork() ){
        case -1:
            oops("cannot fork", 2);

            /* child writes to pipe every 5 seconds */
        case 0:
            len = strlen(CHILD_MESS);
            while ( 1 ){
                if (write(pipefd[1], CHILD_MESS, len) != len )
                    oops("write", 3);
                sleep(5);
            }

            /* parent reads from pipe and also writes to pipe */
        default:
            len = strlen( PAR_MESS );
            while ( 1 ){
                if ( write( pipefd[1], PAR_MESS, len) != len )
                    oops("write", 4);
                sleep(1);
                read_len = read( pipefd[0], buf, BUFSIZ );
                if ( read_len <= 0 )
                    break;
                write( 1, buf, read_len );
            }
        }
    }
}
```

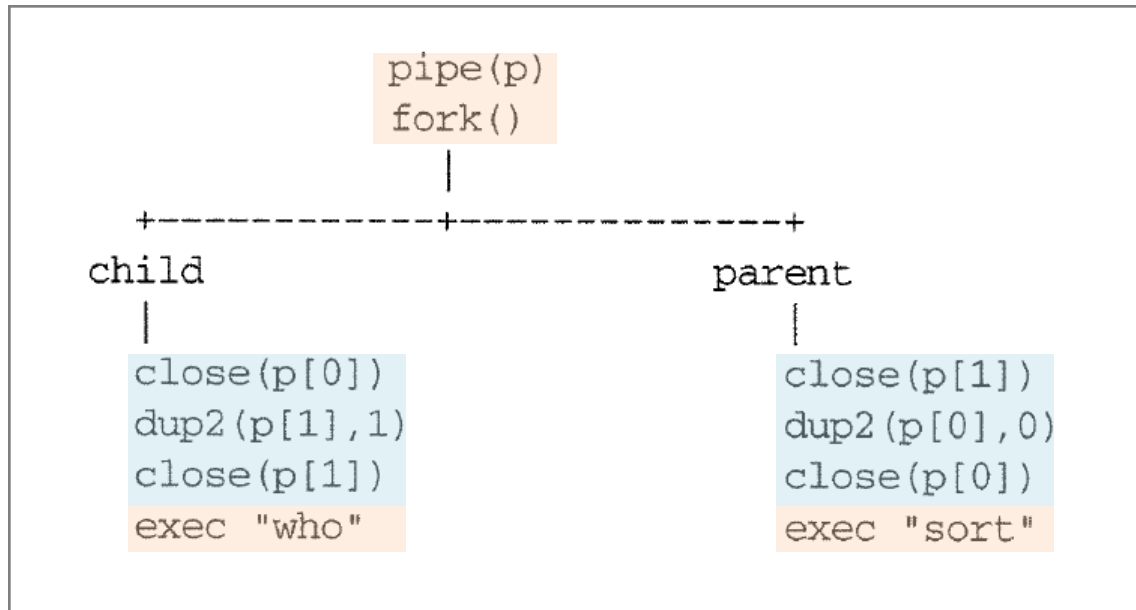
10.6.3 The Finale: Using pipe, fork, and exec

- Writing a general-purpose program `pipe`:

```
$ who | sort
```


```
$ ls | head
```

- Logic



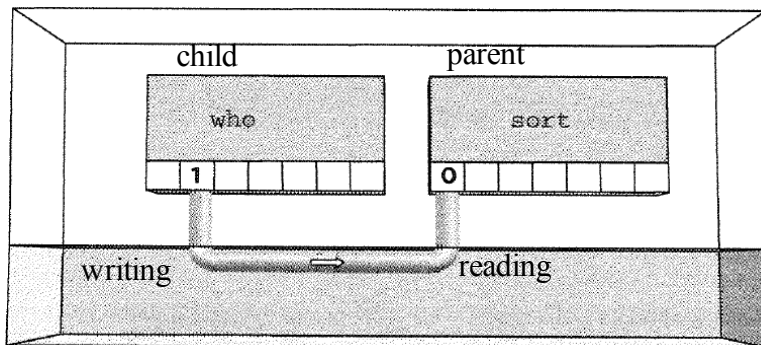
pipe.c

pipe(p)
fork()



p[0] : 3, p[1] : 4

```
+-----+-----+  
child                                     parent  
|  
close(p[0])                             close(p[1])  
dup2(p[1],1)                             dup2(p[0],0)  
close(p[1])                             close(p[0])  
exec "who"                               exec "sort"
```



```
/* pipe.c
 *
 * Demonstrates how to create a pipeline from one process to another
 *
 * Takes two args, each a command, and connects
 *
 * av[1]'s output to input of av[2]
 *
 * usage: pipe command1 command2
 *
 * effect: command1 | command2
 *
 * Limitations: commands do not take arguments
 *
 * uses execlp() since known number of args
 *
 * Note: exchange child and parent and watch fun
 */
#include <stdio.h>
#include <unistd.h>


#define oops(m,x) { perror(m); exit(x); }

main(int ac, char **av)
{
    int    thepipe[2],          /* two file descriptors */
           newfd,               /* useful for pipes      */
           pid;                 /* and the pid           */

    if ( ac != 3 ){
        fprintf(stderr, "usage: pipe cmd1 cmd2\n");
        exit(1);
    }
    if ( pipe( thepipe ) == -1 ) /* get a pipe            */
        oops("Cannot get a pipe", 1);

    /* ----- */
    /* now we have a pipe, now let's get two processes */
    if ( (pid = fork()) == -1 ) /* get a proc           */
        oops("Cannot fork", 2);
```

```
# ./pipe ls head
```



```
/* -----  
/*      Right Here, there are two processes  
/*      parent will read from pipe  
if ( pid > 0 ){                               /* parent will exec av[2] */  
    close(thepipe[1]);                        /* parent doesn't write to pipe */  
    if ( dup2(thepipe[0], 0) == -1 )  
        oops("could not redirect stdin",3);  
    close(thepipe[0]);                        /* stdin is duped, close pipe */  
    execlp( av[2], av[2], NULL);  
    oops(av[2], 4);  
}  
  
/*      child execs av[1] and writes into pipe  
close(thepipe[0]);                            /* child doesn't read from pipe */  
if ( dup2(thepipe[1], 1) == -1 )  
    oops("could not redirect stdout", 4);  
close(thepipe[1]);                            /* stdout is duped, close pipe */  
execlp( av[1], av[1], NULL);  
oops(av[1], 5);  
}
```

```
root@DESKTOP-K4MA2V5:~# ./pipe | ls head
```

```
DESKTOP-K4MA2V5
```

```
a
```

```
a.out
```

```
a.txt
```

```
data
```

```
demodir
```

```
fido.7
```

```
file1.txt
```

```
file2.txt
```

```
file3
```

```
root@DESKTOP-K4MA2V5:~#
```

Technical Details: Pipes Are Not Files

- Pipes were a special case, as they were blocks allocated but not associated to disk blocks.
 - pipe don't have any associated disk blocks associated with it.
 - Linux has a VFS (virtual file system) module called pipefs, that gets mounted in kernel space during boot
 - pipefs cannot be directly examined by the user unlike most file systems
 - The pipe(2) syscall is used by shells and other programs to implement piping, and just creates a new file in pipefs, returning two file descriptors (one for the read end, opening using O_RDONLY, and one for the write end, opened using O_WRONLY)
 - pipefs is stored using an in-memory file system