

# Ch11. Connecting to Processes Near and Far Servers and Sockets

Prof. Seokin Hong

Kyungpook National University

Fall 2018

# Chapter Summary

---

- client/server programming using pipes and sockets
- Interprocess communication and client/server design.
- The idea and techniques of socket programming

# Common Interface for different types of sources

---

## ■ Disk/Device File

- Use `open()` to connect, and use **`read()`** and **`write()`** to transfer data.

## ■ Pipes

- Use `pipe()` to create, use `fork()` to share, and use **`read()`** and **`write()`** to transfer data

## ■ Sockets

- Use `socket()`, `listen()`, and `connect()` to connect, and use **`read()`** and **`write()`** to transfer data.

Use **`read()`** and **`write()`** system call to transfer data to/from processes

# bc: A UNIX CALCULATOR

- Bc can handle very long number:

```
seokin@compasslab1:~$ bc
bc 1.07.1
Copyright 1991-1994, 1997, 1998, 2000, 2004, 2006, 2008, 2012-2017 Free Software Foundation, Inc.
This is free software with ABSOLUTELY NO WARRANTY.
For details type `warranty'.
12^123
54864732218924221509340821667217308113486679281006716245125170218434\
56541709523359082780720277398867836972367369456704108169294512128
1000+1000
2000
20000000+200000
20200000
```

- But bc Is Not a calculator!




FIGURE 11.3

`bc` and `dc` as coroutines.

# bc: A UNIX CALCULATOR

---

## ■ Ideas from bc

### ○ Client/Server Model

- dc provides a service (calculation)
- bc provides a user interface AND uses the service
- bc is called a client of dc.

### ○ Bidirectional Communication

- Using pipes, you need two pipes

### ○ Persistent Service (coroutines)

- Both processes continue to run
- Control passes from one to the other as each completed its part of the job
- bc has the job of parsing input and printing the computation results
- dc has the job of computing




FIGURE 11.3

`bc` and `dc` as coroutines.

# Coding bc: pipe, fork, dup, exec





FIGURE 11.4

`bc`, `dc`, and kernel.

- (a) Create two pipes.
- (b) Create a process to run `dc`.
- (c) In the new process, redirect `stdin` and `stdout` to the pipes, then `exec dc`.
- (d) In the parent, read and parse user input, write commands to `dc`, read response from `dc`, and send response to user.

```

__/**      tinybc.c      * a tiny calculator that uses dc to do its work  __
**
**      * demonstrates bidirectional pipes
**
**      * input looks like number op number which
**      tinybc converts into number \n number \n op \n p
**      and passes result back to stdout
**
**
**      +-----+      +-----+
**      stdin  >0      >== pipetodc ====>      |
**      |      tinybc  |      |      dc -      |
**      stdout <1      <== pipefromdc ==<      |
**      +-----+      +-----+
**
**
**      * program outline
**      a. get two pipes
**      b. fork (get another process)
**      c. in the dc-to-be process,
**          connect stdin and out to pipes
**          then execl dc
**      d. in the tinybc-process, no plumbing to do
**          just talk to human via normal i/o
**          and send stuff via pipe
**      e. then close pipe and dc dies
**      * note: does not handle multiline answers
**/

```





---


```

void be_dc(int in[2], int out[2])
/*
 *      set up stdin and stdout, then execl dc
 */
{
    /* setup stdin from pipein */
    if ( dup2(in[0],0) == -1 )      /* copy read end to 0 */
        oops("dc: cannot redirect stdin",3);
    close(in[0]);                  /* moved to fd 0 */
    close(in[1]);                  /* won't write here */


    /* setup stdout to pipeout */
    if ( dup2(out[1], 1) == -1 )    /* dupe write end to 1 */
        oops("dc: cannot redirect stdout",4);
    close(out[1]);                  /* moved to fd 1 */
    close(out[0]);                  /* won't read from here */

    /* now execl dc with the - option */
    execlp("dc", "dc", "-", NULL );
    oops("Cannot run dc", 5);
}

```



Standard input stream will be processed.




```
void be_bc(int todc[2], int fromdc[2])
/*
 *   read from stdin and convert into to RPN, send down pipe
 *   then read from other pipe and print to user
 *   Uses fdopen() to convert a file descriptor to a stream
 */
{
    int    num1, num2;
    char    operation[BUFSIZ], message[BUFSIZ], *fgets();
    FILE    *fpout, *fpin, *fdopen();

    /* setup */
    close(todc[0]);           /* won't read from pipe to dc */
    close(fromdc[1]);         /* won't write to pipe from dc */ ④

    fpout = fdopen( todc[1], "w" );      /* convert file desc- */
    fpin  = fdopen( fromdc[0], "r" );     /* riptors to streams */ ⑤
    if ( fpout == NULL || fpin == NULL )
        fatal("Error converting pipes to streams");
}
```

---

```
/* main loop */
while ( printf("tinybc: "), fgets(message, BUFSIZ, stdin) != NULL ) {
    /* parse input */
    if ( sscanf(message, "%d%[-+*/^]%d", &num1, operation,
        &num2) != 3 ) {
        printf("syntax error\n");
        continue;
    }

    if ( fprintf( fpout , "%d\n%d\n%c\n", num1, num2,
        *operation ) == EOF )  2\n2\n+\n
        fatal("Error writing");
    fflush( fpout );
    if ( fgets( message, BUFSIZ, fpin ) == NULL )
        break;
    printf("%d %c %d = %s", num1, *operation, num2, message); // stdout
}
fclose(fpout);          /* .close pipe          */
fclose(fpin);           /* dc will see EOF          */
}
```

---

```
void fatal(char mess[])
{
    fprintf(stderr, "Error: %s\n", mess);
    exit(1);
}
```

```
$ cc tinybc.c -o tinybc ; ./tinybc
```

```
tinybc: 2+2 → no spaces
```

```
2 + 2 = 4
```

```
tinybc: 55^5
```

```
55 ^ 5 = 503284375
```

```
tinybc:
```

```
...
```

```
tinybc: ctrl+D
```

```
$
```

# fdopen: Making File Descriptors Look like Files

---

- `fopen`: file name → `FILE *`
- `fdopen`: file descriptor → `FILE *`
  - you can use **standard, buffered I/O** operations;
  - **`tinybc.c` uses `fprintf` and `fgets` to send data through the pipes to dc.**

# Contents

---

- bc: A Unix Calculator
- popen: Making Process Look like Files
- Sockets: Connecting to Remote Processes

# popen: MAKING PROCESSES LOOK LIKE FILE

---

- We examine the popen library function.
- We see what popen does and how popen works, and then we write our own version.

# What popen Does

---

- `fopen()` opens a buffered connection to a file:

```
FILE *fp;                                /* a pointer to a struct */
fp = fopen( "file1", "r" );              /* args are filename, connection type */
c = getc(fp);                            /* read char by char */
fgets(buf, len, fp);                     /* line by line */
fscanf(fp, "%d%d%s", &x, &y, x);          /* token by token */
fclose(fp);                              /* close when done */
```

- `popen()` opens a buffered connection to a process:

```
FILE *fp;                                /* same type of struct */
fp = popen("ls", "r");                    /* args are program name, connection type */
fgets(buf, len, fp);                     /* exactly the same functions */
pclose(fp);                              /* close when done */
```



# What popen Does

---

- Similarities between popen and fopen.

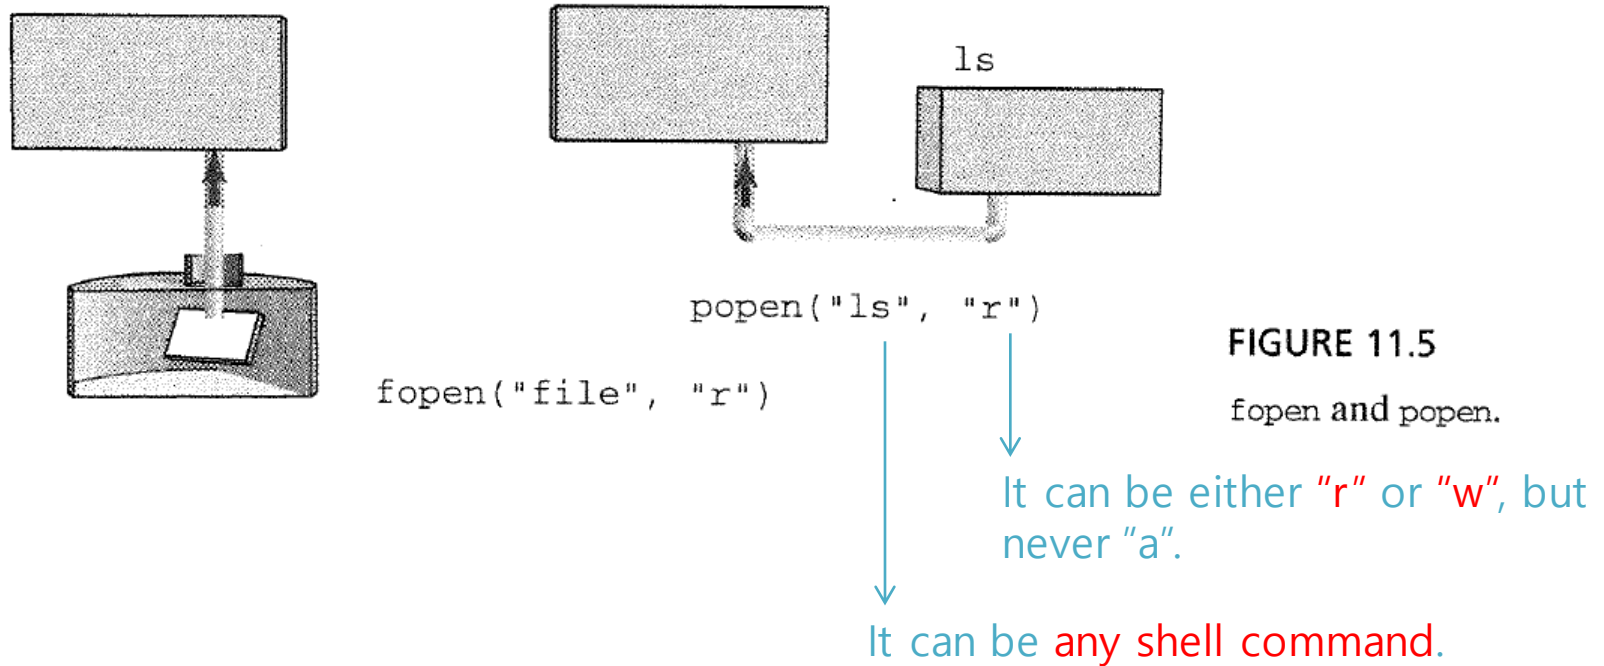


FIGURE 11.5

`fopen` and `popen`.

```
/* popendemo.c  
 * demonstrates how to open a program for standard i/o  
 * important points:  
 *      1. popen() returns a FILE *, just like fopen()  
 *      2. the FILE * it returns can be read/written  
 *         with all the standard functions  
 *      3. you need to use pclose() when done  
 */  
  
#include <stdio.h>  
#include <stdlib.h>  
  
int main()  
{  
    FILE *fp;  
    char buf[100];  
    int i = 0;  
  
    fp = popen( "who | sort", "r" );          /* open the command */  
    while ( fgets( buf, 100, fp ) != NULL ) /* read from command */  
        printf("%3d %s", i++, buf );         /* print data       */  
    pclose( fp );                             /* IMPORTANT!           */  
    return 0;  
}
```

---

- `pclose` is Required

- A process needs to be waited for, or it becomes a *zombie*.
- `pclose` calls `wait`.

# Writing popen: Using fdopen

---

- How does `popen` work?
  - `popen` **runs** a program and **returns** a connection to the standard input or standard output of that program.

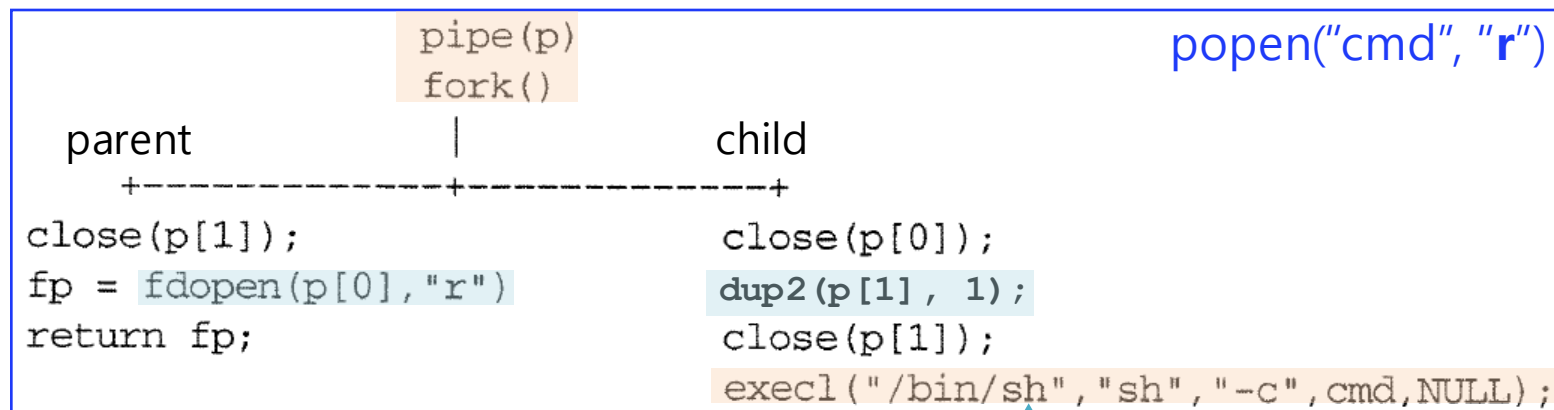
```
FILE    *fp;
char    buf[100];
int      i = 0;

fp = popen( "who|sort", "r" );

while ( fgets( buf, 100, fp ) != NULL )
    printf("%3d %s", i++, buf );

pclose( fp );
return 0;
```

## ■ Writing popen: Using fdopen



```
FILE *fp;  
char buf[BUFSIZ];  
  
fp = popen("ls", "r");  
while( fgets(buf, BUFSIZ, fp) != NULL )  
    fputs(buf, stdout);  
return 0;
```

The only program that can run any shell command is the shell itself




FIGURE 11.6

Reading from a shell command.

---

```
/* popen.c - a version of the Unix popen() library function
 *      FILE *popen( char *command, char *mode )
 *          command is a regular shell command
 *          mode is "r" or "w"
 *          returns a stream attached to the command, or NULL
 *          execl's "sh" "-c" command
 *      todo: what about signal handling for child process?
 */
#include <stdio.h>
#include <signal.h>
#define READ 0
#define WRITE 1

const char* given_command = "who | sort";
const char* given_mode = "r";

void main(){

    FILE* fp = popen(given_command,given_mode);
    char buf[100];
    int i=0;

    while(fgets(buf,100,fp)!=NULL)
        printf("%3d %s", i++, buf);
    pclose(fp);

    return 0;
}
```

---

```
FILE *popen(const char *command, const char *mode)
{
    int      pfp[2], pid;                /* the pipe and the process */
    FILE      *fdopen(), *fp;            /* fdopen makes a fd a stream */
    int      parent_end, child_end;      /* of pipe */
    if ( *mode == 'r' ){                 /* figure out direction */
        parent_end = READ;
        child_end = WRITE ;
    } else if ( *mode == 'w' ){
        parent_end = WRITE;
        child_end = READ ;
    } else return NULL ;

    if ( pipe(pfp) == -1 )                /* get a pipe */
        return NULL;
    if ( (pid = fork()) == -1 ){          /* and a process */
        close(pfp[0]);                   /* or dispose of pipe */
        close(pfp[1]);
        return NULL;
    }
}
```

---

```
/* ----- parent code here ----- */
/*   need to close one end and fdopen other end   */

if ( pid > 0 ){
    if (close( pfp[child_end] ) == -1 )
        return NULL;
    return fdopen( pfp[parent_end] , mode );    /* same mode */
}

/* ----- child code here ----- */
/*   need to redirect stdin or stdout then exec the cmd   */

if ( close(pfp[parent_end]) == -1 )    /* close the other end */
    exit(1);                          /* do NOT return      */

if ( dup2(pfp[child_end], child_end) == -1 )
    exit(1);

if ( close(pfp[child_end]) == -1 )    /* done with this one */
    exit(1);

                                   /* all set to run cmd   */
execl( "/bin/sh", "sh", "-c", command, NULL );
exit(1);

}
```



# Contents


---

- bc: A Unix Calculator
- popen: Making Process Look like Files
- Sockets: Connecting to Remote Processes

# Sockets

---

- Limitations of pipes
  - It can only connect related processes : ex) parent/child
  - It can only connect processes on the same computer
- Unix provides another method of inter-process communication : **sockets**
- Sockets allow processes to create pipe-like connections to unrelated processes and even to process on other computers



# Important Concepts in Socket Programming

---

## ■ Client and server

- In Unix terms, server is a program, not a computer
- A server process wait for a request processes it and loops back to take the next request but a client does not loop.

## ■ Hostname and port:

- A server on the internet is a process running on a computer
- The computer is called the host, and it has a hostname and a port number.

## ■ Address family

- Like phone-number or zip code, a host has an address.

## ■ Protocol

- A set of rules that governs interaction between client and server.

# Example: Time Server and Client




FIGURE 11.9


Processes on different computers.

```
seokin@compasslab1:~/workspace/system_programming/labs/lab10$ ./timeserv &  
[1] 21714  
seokin@compasslab1:~/workspace/system_programming/labs/lab10$ ./timeclnt  
compasslab1 23000  
Wow! got a call!!!  
The time here is ..Tue Nov 13 23:49:31 2018  
seokin@compasslab1:~/workspace/system_programming/labs/lab10$
```

# Writing timeserv.c: A Time Server

---

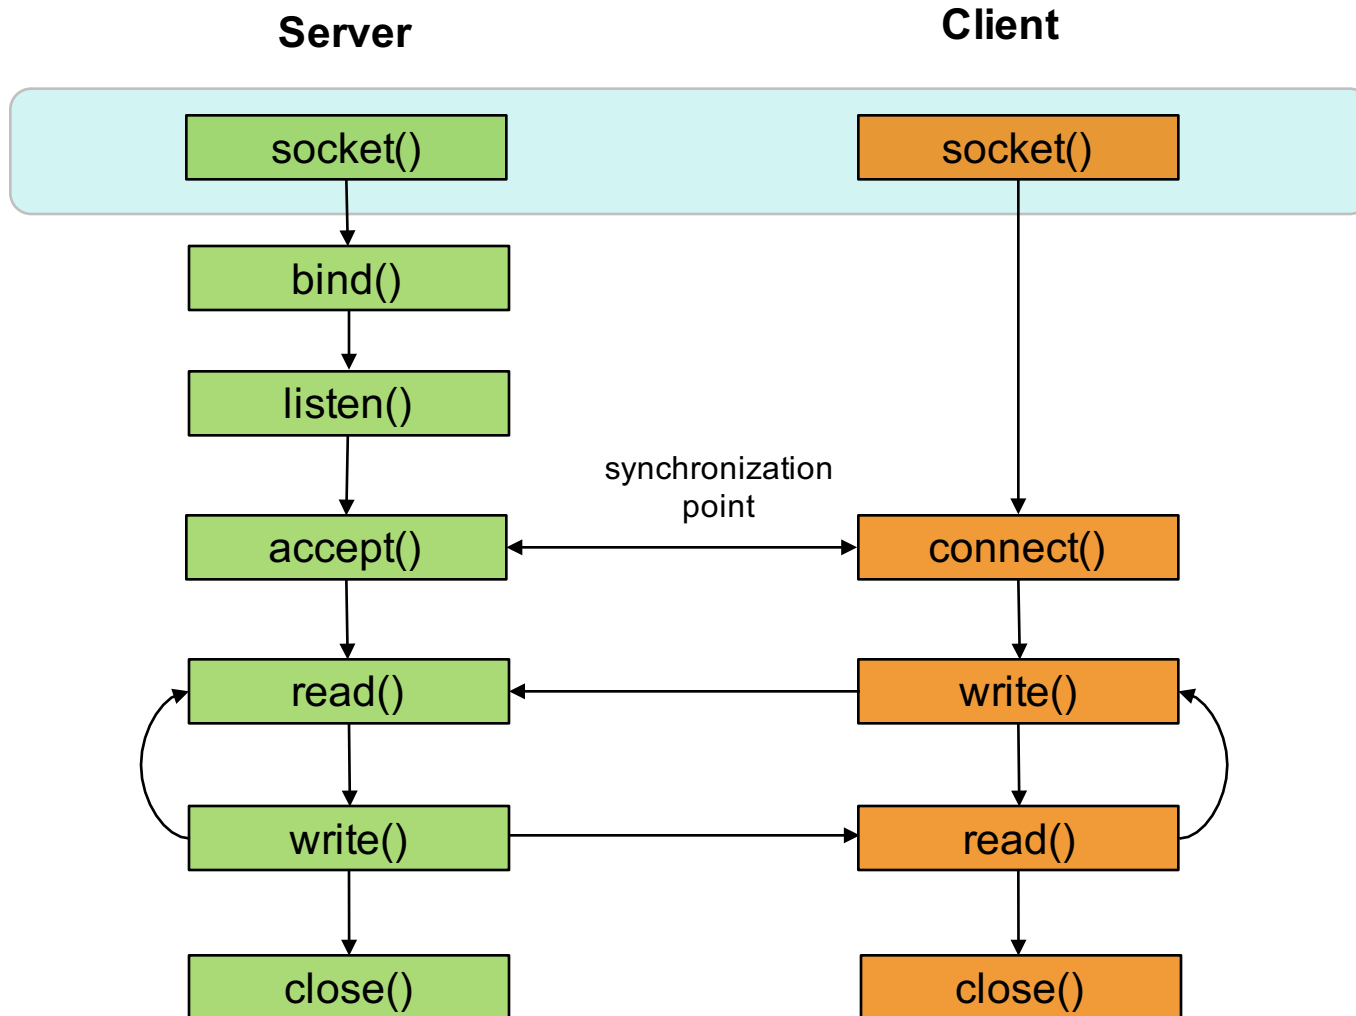
action	syscall
1. Get a phone line	socket
2. Assign a number	bind
3. Allow incoming calls	listen
4. Wait for a call	accept
5. Transfer data	read/write
6. Hang up	close



# Writing timeserv.c: A Time Server

---

- Step 1: Ask kernel for a socket



# Writing timeserv.c: A Time Server

---

## ■ Step 1: Ask kernel for a socket

---

### socket

---

PURPOSE	Create a socket		
INCLUDE	#include <sys/types.h> #include <sys/socket.h>		
USAGE	sockid = socket(int domain, int type, int protocol)		
ARGS	domain	communication domain. PF_INET is for Internet sockets	PF_INET : IPv4 protocols (internet addresses) PF_UNIX : local communication
	type	type of socket. SOCK_STREAM looks like a pipe	SOCK_STREAM : reliable, 2-way, connection-based service SOCK_DGRAM : unreliable, connectionless
	protocol	protocol used within the socket. 0 is default.	IPPROTO_TCP, IPPROTO_UDP 0: system will determine the protocol
RETURNS	-1	if error	
	sockid	a socket id if successful	

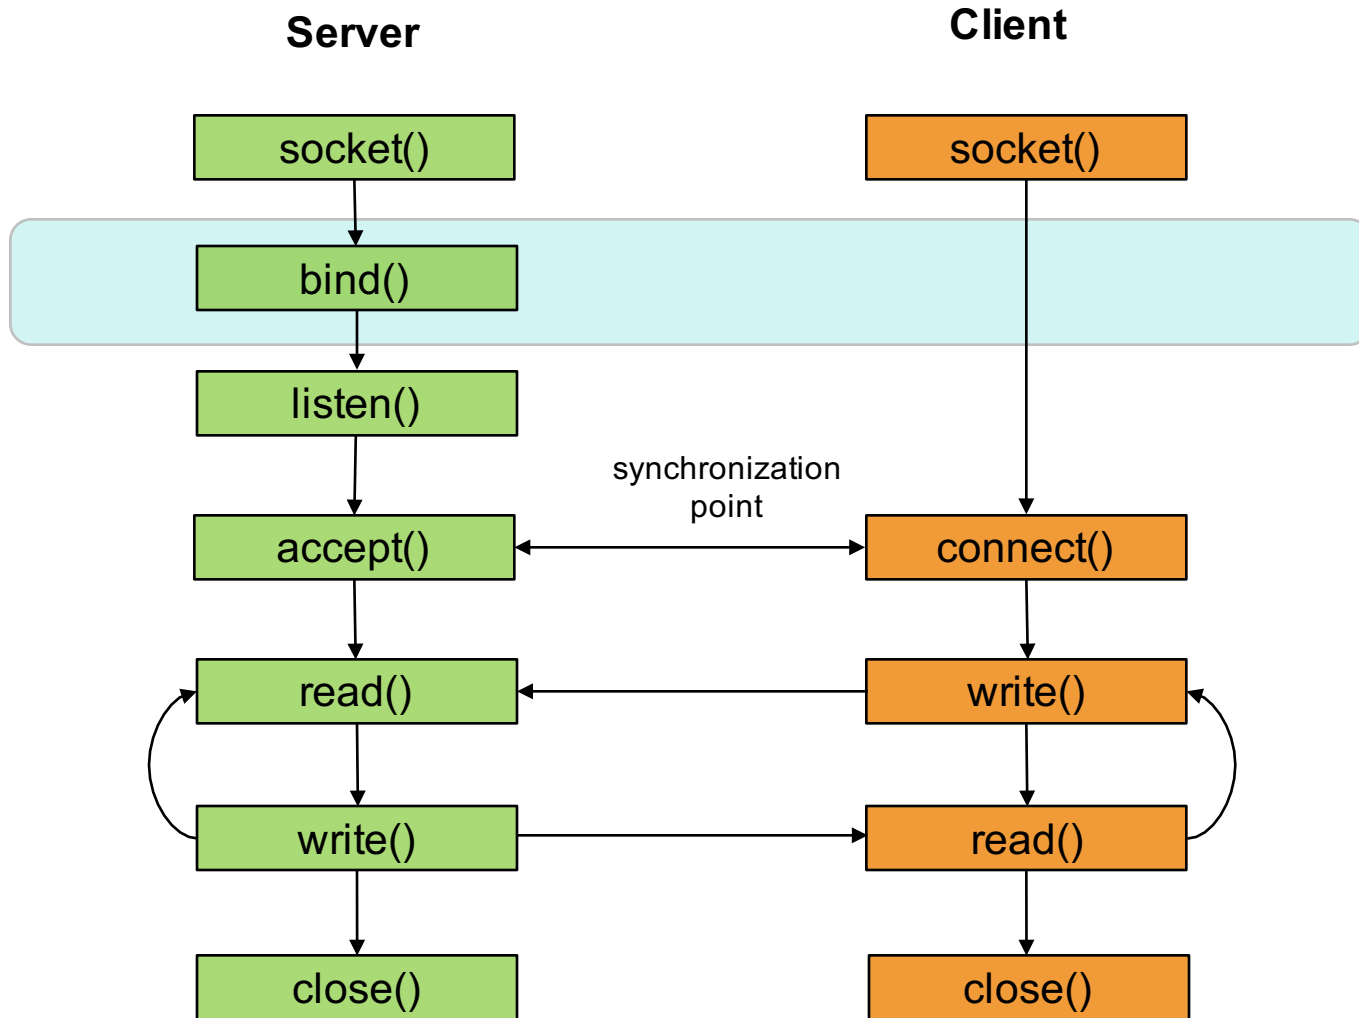
---

```
sock_id = socket( PF_INET, SOCK_STREAM, 0 );
if ( sock_id == -1 )
    oops( "socket" );
```

# Writing `timeserv.c`: A Time Server

---

- Step 2: Bind address to socket; Address is host, port.





# Writing timeserv.c: A Time Server

---

- Step 2: Bind address to socket; Address is host, port.

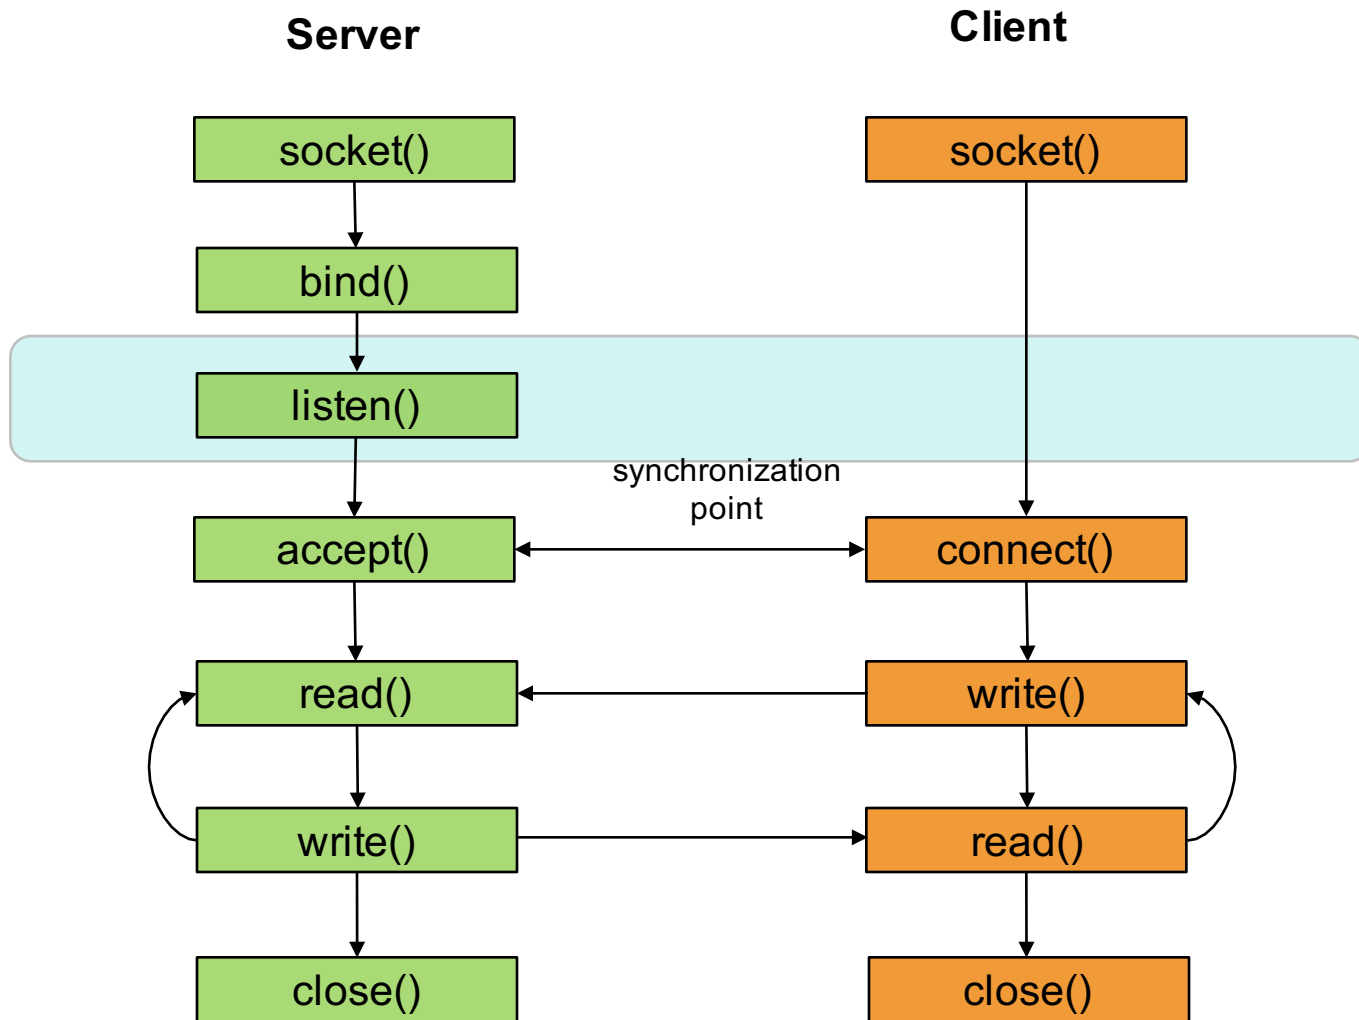
bind		
PURPOSE	Bind an address to a socket	
INCLUDE	#include <sys/types.h> #include <sys/socket.h>	
USAGE	result = bind(int sockid, struct sockaddr *addrp, socklen_t addrlen)	
ARGS	sockid	the id of the socket
	addrp	a pointer to a struct containing the address
	addrlen	the length of the struct
RETURNS	-1	if error
	0	if success

```
if ( bind(sock_id, (struct sockaddr *)&saddr, sizeof(saddr)) != 0 )  
    oops( "bind" );
```

# Writing `timeserv.c`: A Time Server

---

- Step 3: Allow incoming calls on socket



# Writing timeserv.c: A Time Server

---

## ■ Step 3: Allow incoming calls on socket

---

### **listen**

---

<b>PURPOSE</b>	Listen for connections on a socket	
<b>INCLUDE</b>	#include <sys/socket.h>	
<b>USAGE</b>	result = listen(int sockid, int qsize)	
<b>ARGS</b>	sockid	socket that will accept calls
	qsize	backlog of incoming connections
<b>RETURNS</b>	-1	if error
	0	if success

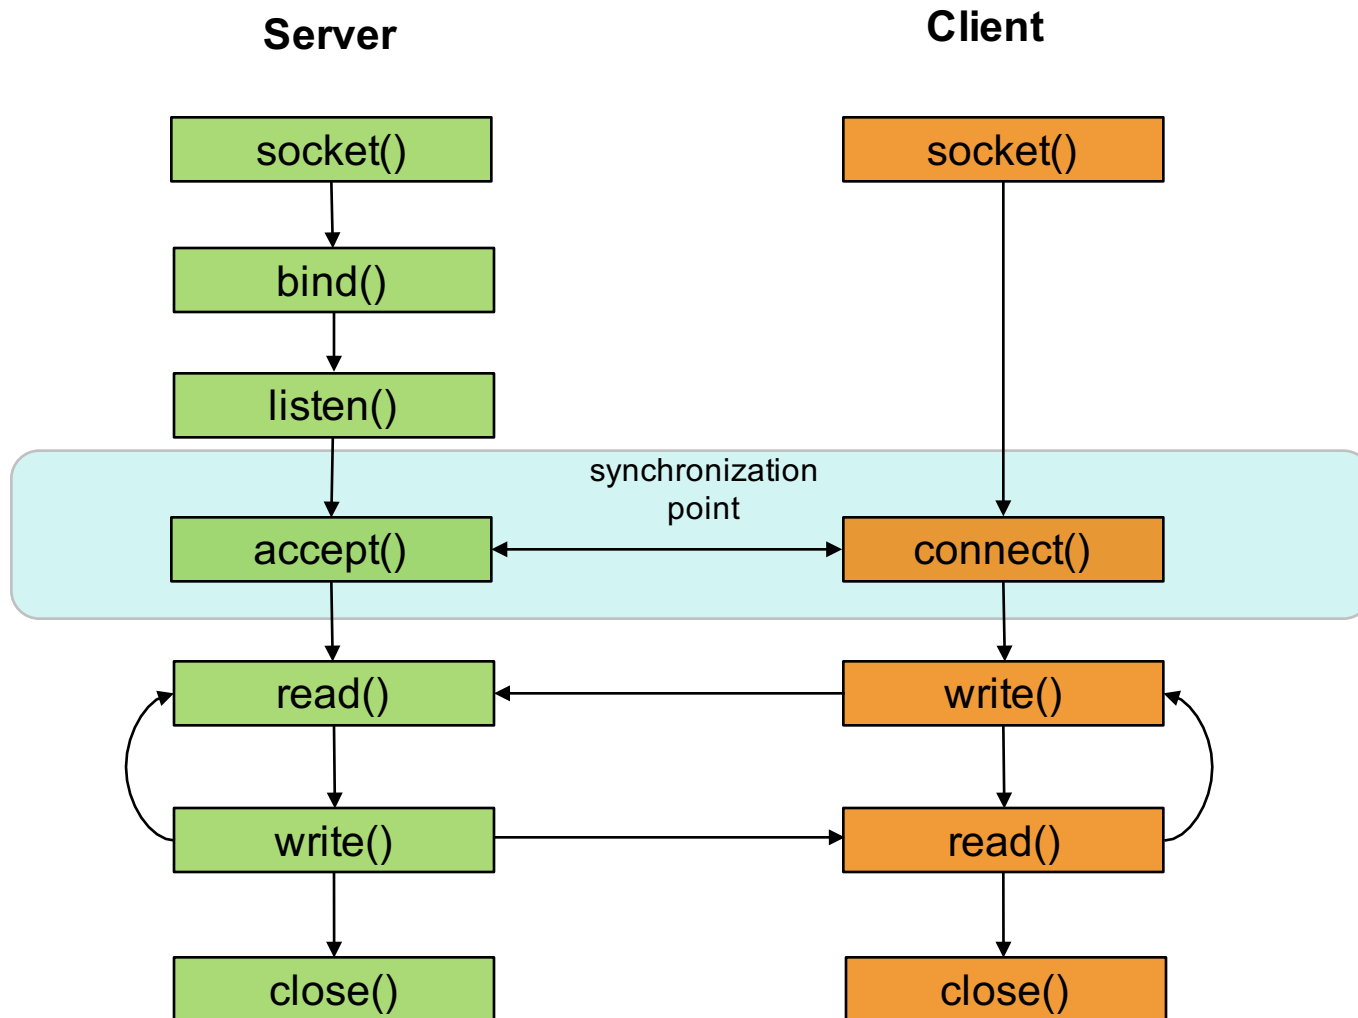
---

```
if ( listen(sock_id, 1) != 0 )  
    oops( "listen" );
```

# Writing timeserv.c: A Time Server

---

## ■ Step 4: Wait For/Accept a connection



# Writing timeserv.c: A Time Server

---

## ■ Step 4: Wait For/Accept a connection

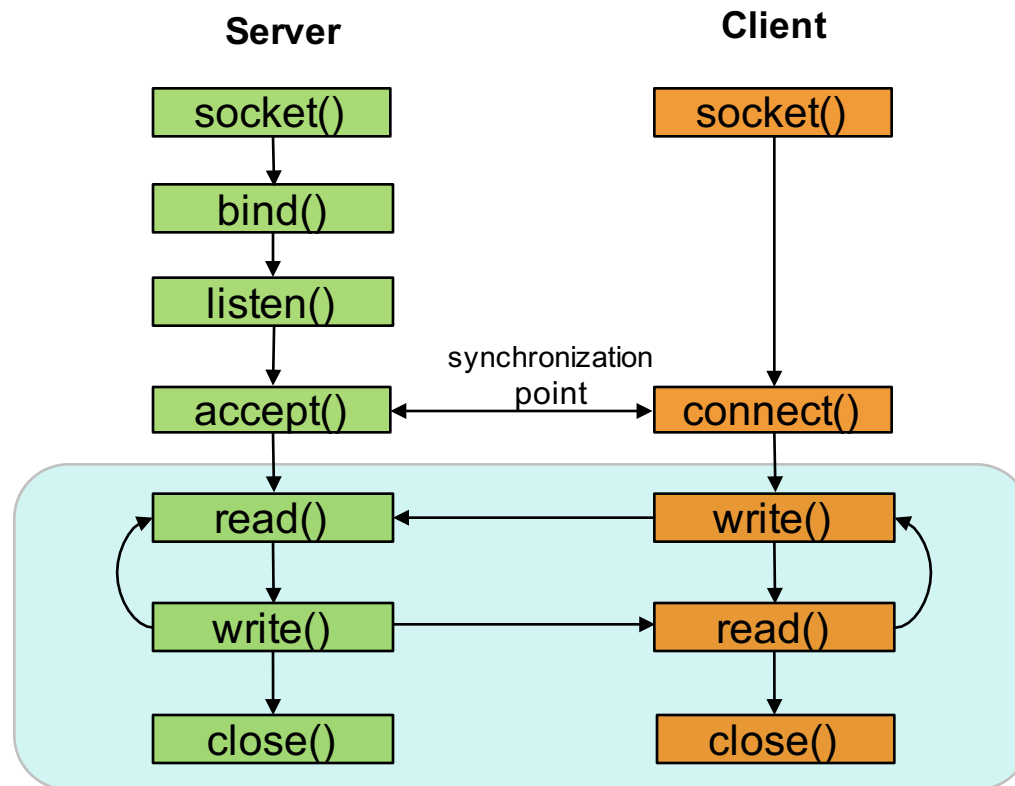
accept		
PURPOSE	Accept a connection on a socket	
INCLUDE	#include <sys/types.h> #include <sys/socket.h>	
USAGE	fd = accept(int sockid, struct sockaddr *callerid, socklen_t *addrlenp)	
ARGS	sockid	accept a call on this socket
	callerid	pointer to struct for address of caller
	addrlenp	pointer to storage for length of address of caller
RETURNS	-1	if error
	fd	a file descriptor open for reading and writing

```
sock_fd = accept(sock_id, NULL, NULL);
```

# Writing `timeserv.c`: A Time Server

---

- Step 5: Transfer Data
- Step 6: Close Connection



# Writing timeserv.c: A Time Server

---

```
/* timeserv.c - a socket-based time of day server
 */
#include <stdlib.h>          //for exit()
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <time.h>
#include <strings.h>
#define PORTNUM 13000      /* our time service phone number */
#define HOSTLEN 256
#define oops(msg)          { perror(msg) ; exit(1) ; }
```

---

```
int main(int ac, char *av[])
{
    struct sockaddr_in  saddr;    /* build our address here */
    struct hostent      *hp;      /* this is part of our    */
    char                hostname[HOSTLEN]; /* address                */
    int                 sock_id, sock_fd; /* line id, file desc     */
    FILE                *sock_fp;   /* use socket as stream   */
    char                *ctime();    /* convert secs to string */
    time_t              thetime;     /* the time we report     */

    /*
     * Step 1: ask kernel for a socket
     */

    sock_id = socket( PF_INET, SOCK_STREAM, 0 ); /* get a socket */
    if ( sock_id == -1 )
        oops( "socket" );
```



```
/*  
 * Step 2: bind address to socket.  Address is host,port  
 */  
  
bzero( (void *)&saddr, sizeof(saddr) ); /* clear out struct */  
  
gethostname( hostname, HOSTLEN );      /* where am I ? */  
hp = gethostbyname( hostname );        /* get info about host */  
                                         /* fill in host part */  
bcopy( (void *)hp->h_addr, (void *)&saddr.sin_addr, hp->h_length);  
saddr.sin_port = htons(PORTNUM);      /* fill in socket port */  
saddr.sin_family = AF_INET ;          /* fill in addr family */  
  
if ( bind(sock_id, (struct sockaddr *)&saddr, sizeof(saddr)) != 0 )  
    oops( "bind" );  
  
/*  
 * Step 3: allow incoming calls with Qsize=1 on socket  
 */  
  
if ( listen(sock_id, 1) != 0 )  
    oops( "listen" );
```

\* htons(): translate a short integer from host byte order to network byte order

---

```
/*
 * main loop: accept(), write(), close()
 */

while ( 1 ){
    sock_fd = accept(sock_id, NULL, NULL); /* wait for call */
    printf("Wow! got a call!\n");
    if ( sock_fd == -1 )
        oops( "accept" );          /* error getting calls */

    sock_fp = fdopen(sock_fd, "w"); /* we'll write to the */
    if ( sock_fp == NULL )          /* socket as a stream */
        oops( "fdopen" );          /* unless we can't */

    thetime = time(NULL);           /* get time */
                                    /* and convert to string */
    fprintf( sock_fp, "The time here is .." );
    fprintf( sock_fp, "%s", ctime(&thetime) );
    fclose( sock_fp );              /* release connection */
}
}
```

---

```
struct hostent {
    char*h_name; /* official name of host */
    char**h_aliases; /* alias list */
    int h_addrtype; /* host address type */
    int h_length; /* length of address */
    char**h_addr_list; /* list of addresses */
};
#define h_addr h_addr_list[0] /* for backward compatibility */
```

```
struct sockaddr {
    unsigned short sa_family; /* Address family (e.g. AF_INET) */
    char sa_data[14]; /* Family-specific address information */
}
```


```
struct in_addr {
    unsigned long s_addr; /* Internet address (32 bits) */
}
```

```
struct sockaddr_in {
    unsigned short sin_family; /* Internet protocol (AF_INET) */
    unsigned short sin_port; /* Address port (16 bits) */
    struct in_addr sin_addr; /* Internet address (32 bits) */
    char sin_zero[8]; /* Not used */
}
```

# Writing timeclnt.c: a Time Client

---

action	syscall
1. Get a phone line	socket
2. Call the server	connect
3. Transfer data	read/write
4. Hang up	close



# Writing timeclnt.c: a Time Client

---

## ■ Establish a connection

<b>connect</b>		
<b>PURPOSE</b>	Connect to a socket	
<b>INCLUDE</b>	#include <sys/types.h> #include <sys/socket.h>	
<b>USAGE</b>	result = connect(int sockid, struct sockaddr *serv_addrp, socklen_t addrlen);	
<b>ARGS</b>	sockid	socket to use for connection
	serv_addrp	pointer to struct containing server address
	addrlen	length of that struct
<b>RETURNS</b>	-1	if error
	0	if success

---

```
/* timeclnt.c - a client for timeserv.c
 *          usage: timeclnt hostname portnumber
 */
#include      <stdio.h>
#include      <sys/types.h>
#include      <sys/socket.h>
#include      <netinet/in.h>
#include      <netdb.h>
#include      <stdlib.h>
#include      <strings.h>

#define      oops(msg)      { perror(msg); exit(1); }

main(int ac, char *av[])
{
    struct sockaddr_in servadd;      /* the number to call */
    struct hostent      *hp;          /* used to get number */
    int      sock_id, sock_fd;        /* the socket and fd */
    char      message[BUFSIZ];        /* to receive message */
    int      messlen;                 /* for message length */
```

```

/*
 * Step 1: Get a socket
 */

sock_id = socket( AF_INET, SOCK_STREAM, 0 );    /* get a line */
if ( sock_id == -1 )
    oops( "socket" );                          /* or fail */

/*
 * Step 2: connect to server
 *      need to build address (host,port) of server first
 */

bzero( &servadd, sizeof( servadd ) );    /* zero the address */
hp = gethostbyname( av[1] );              /* lookup host's ip # */
if (hp == NULL)
    oops(av[1]);                          /* or die */
bcopy(hp->h_addr, (struct sockaddr *)&servadd.sin_addr, hp->h_length);
servadd.sin_port = htons(atoi(av[2]));  /* fill in port number */
servadd.sin_family = AF_INET ;           /* fill in socket type */

                                           /* now dial */
if ( connect(sock_id,(struct sockaddr *)&servadd, sizeof(servadd)) !=0)
    oops( "connect" );

```

```
/*  
 * Step 3: transfer data from server, then hangup  
 */  
  
    messlen = read(sock_id, message, BUFSIZ);      /* read stuff   */  
    if ( messlen == - 1 )  
        oops("read") ;  
  
    if ( write( 1, message, messlen ) != messlen ) /* and write to */  
        oops( "write" );                          /* stdout       */  
    close( sock_id );  
}
```