

# Ch12. Connections and Protocols: Writing a Web Server

Prof. Seokin Hong

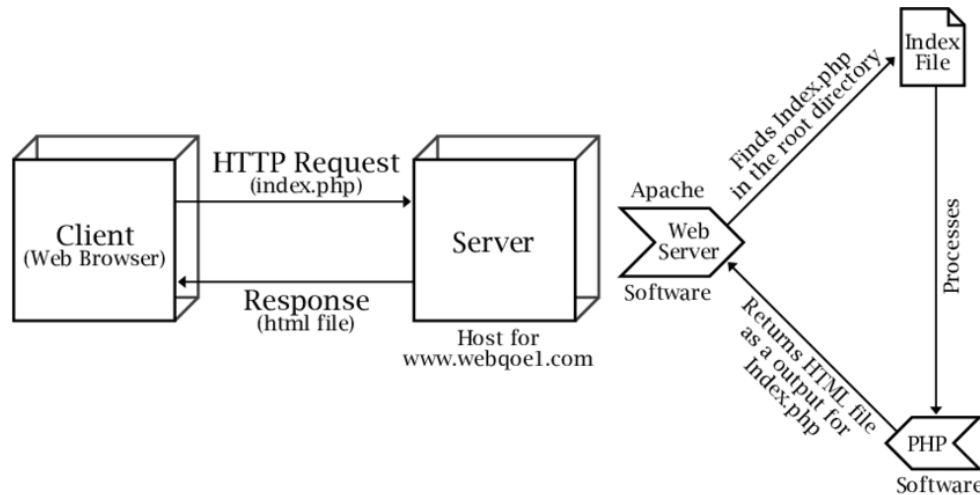
Kyungpook National University

Fall 2018

# Web Server?

---

- Using the world wide web is easy
  - Type a web location into a browser or click on a link, and a web page is delivered from a **remote computer**

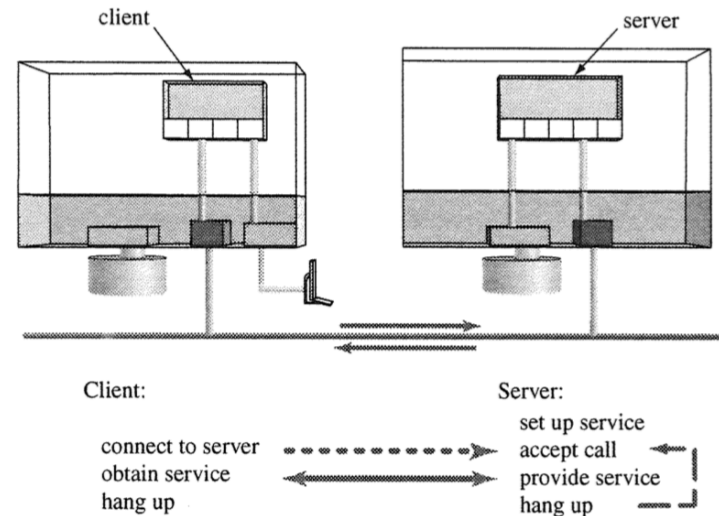


# Three main operations of Server/Client System

---

## ■ Three main operations

- Server sets up a service
- Client connects to a server
- Server and client do business



- Most socket-based client/server systems are pretty similar
  - E-mail, file transfer, distributed databases,....
- Once you understand one socket-based client/server system, you can understand most of the other ones

# Setting Up a Server Socket

---

## ■ Three steps

1. Create socket  
`sock = socket(PF_INET, SOCK_STREAM, 0)`
2. Give the socket an address  
`bind(sock, &addr, sizeof(addr))`
3. Arrange to take incoming calls  
`listen(sock, queue_size)`

## ■ The three-step can be combined into a single function

- `sock = make_server_socket (int portnum)`  
return -1 if error, or a server socket listening at port “portnum”

## Operation 2 : Connecting to a Server

---

1. Create a socket  
sock = **socket**(PF\_INET, SOCK\_STREAM, 0)
  2. Use the socket to connect to a server  
**connect**(sock, &serv\_addr, size(serv\_addr))
- 
- The parameters are “hostname” and “port number” of serv\_addr.
  - Thus we can be combined into a single function
    - fd = **connect\_to\_server**(hostname, portnum)  
returns -1 if error, or a fd open for reading and writing  
connected to the socket at port “portnum” on host “hostname”

# socklib.c

---

```
// socklib.h
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <time.h>
#include <strings.h>

// register functions to be implemented
int make_server_socket(int);
int make_server_socket_q(int, int);
int connect_to_server(char*, int);
#define HOSTLEN 256
#define BACKLOG 1 // only one request at a time

int make_server_socket(int portnum){
    return make_server_socket_q(portnum, BACKLOG);
}
```

# socklib.c

---

```
int make_server_socket_q(int portnum, int backlog)
{
    struct sockaddr_in  saddr;          /* build our address here */
    struct hostent*  hp;                /* this is part of our      */
    char  hostname[HOSTLEN];           /* address                  */
    int sock_id;                        /* the socket               */

    sock_id = socket(PF_INET, SOCK_STREAM, 0); /* get a socket */
    if ( sock_id == -1 )
        return -1;

    /** build address and bind it to socket **/

    bzero((void *)&saddr, sizeof(saddr));    /* clear out struct */
    gethostname(hostname, HOSTLEN);            /* where am I ?     */
    hp = gethostbyname(hostname);              /* get info about host */
}
```

# socklib.c

---

```
    bcopy( (void *)hp->h_addr, (void *)&saddr.sin_addr, hp->h_length); /* fill in host part */
    saddr.sin_port = htons(portnum); /* fill in socket port */
    saddr.sin_family = AF_INET ; /* fill in addr family */
    if ( bind(sock_id, (struct sockaddr *)&saddr, sizeof(saddr)) != 0 )
        return -1;

    /** arrange for incoming calls **/
    if ( listen(sock_id, backlog) != 0 )
        return -1;
    return sock_id;
}
```



# socklib.c

---

```
int connect_to_server(char *host, int portnum)
{
    int                sock;
    struct sockaddr_in servadd;    /* the number to call */
    struct hostent      *hp;      /* used to get number */

    /** Step 1: Get a socket **/

    sock = socket( AF_INET, SOCK_STREAM, 0 ); /* get a line */
    if ( sock == -1 )
        return -1;

    /** Step 2: connect to server **/

    bzero( &servadd, sizeof(servadd) );      /* zero the address */
    hp = gethostbyname( host );    /* lookup host's ip # */
    if (hp == NULL)
        return -1;
```

# socklib.c

---

```
bcopy(hp->h_addr, (struct sockaddr*)&servadd.sin_addr, hp->h_length);
servadd.sin_port = htons(portnum);          /* fill in port number */
servadd.sin_family = AF_INET ;              /* fill in socket type */

if ( connect(sock,(struct sockaddr*)&servadd, sizeof(servadd)) !=0)
    return -1;

return sock;
}
```



# timeserv/timecInt using socklib.c

---

- What we need to implement
  - **talk\_with\_server()** for client
  - **process\_request()** for server

```
talk_with_server(fd)
{
    char buf[LEN];
    int  n;

    n=read(fd,buf,LEN);
    write(1,buf,n);
}
```

```
process_request(fd)
{
    time_t now;
    char *cp;

    time(&now);
    cp = ctime(&now);
    write(fd, cp, strlen(cp));
}
```

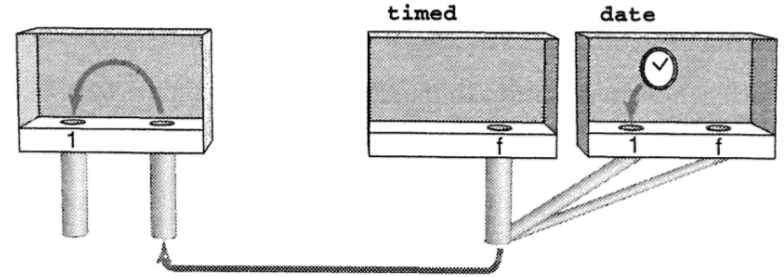
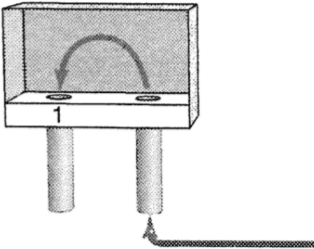
## A Second Version of the Server: Using fork

```
process_request(fd)
/* send the date out to the client via fd */
{
    int pid = fork();

    switch(pid) {
        case -1: return;          /* cannot provide service */

        case 0: dup2(fd, 1);      /* child runs date          */
                close(fd);        /* by redirecting stdout */
                execl("/bin/date", "date", NULL);
                oops("execlp"); /* or quits */

        default: wait(NULL);      /* parent wait for child */
    }
}
```



## What is advantage of the second version of time server?

# Server Design Question: DIY or Delegate

---

- Two types of server design
  - **Do it yourself** : Serer does the work itself
  - **Delegate** : the server forks a process to do the work
  
- **Advantages and disadvantages of each design**
  - DIY for quick, simple tasks
  
  - **Delegate for slower, more complex tasks.**
    - **Can handle many requests simultaneously.**
    - In order to serve several requests at once, a server should not “wait” for child to finish.
    - How to prevent child processes from being zombies?

# How to prevent child processes from being zombies?

---

- Rather than wait for child to die, a parent can arrange to receive a signal when a child dies
- When a child process exists or is killed, kernel sends SIGCHLD (SIGCHLD is ignored by default)
- Parent process may set a signal handler. That handler can call wait.

# Using SIGCHLD for Zombie Prevention

---

```
main(){
    int sock, fd;

    signal(SIGCHLD, child_waiter);
    if((sock = make_server_socket(PORTNUM)) == -1) oops("make_server_socket");
    while(1) {
        fd = accept(sock, NULL, NULL);
        if(fd == -1) break;
        process_request(fd);
        close(fd);
    }
}

void child_waiter(int signum)
{
    wait(NULL);
}
```



# Using SIGCHLD for Zombie Prevention

---

## ■ Problem 1.

- Jump to the signal handler interrupts the accept system call.
- When interrupted by signal, accept returns -1, and sets errno to EINTR.
- Our code treats the value of -1 from accept as an error and breaks from the main loop.
- How to distinguish between real error and an interrupted system call?

# Using SIGCHLD for Zombie Prevention

---

## ■ Problem 2.

- What happens if several child processes exit at almost exactly the same time?
- Unix blocks signals, but does not queue signals. So newly-coming signals are lost.
- If children exit while the parent is in the handler, signals from the children are lost.
- How to prevent this?

# Using SIGCHLD for Zombie Prevention

---

## ■ Solution

Call wait enough times to mop up all terminated processes.

```
void child_waiter (int signum)
{
    while (waitpid (-1, NULL, WNOHANG) > 0) ;
}
```

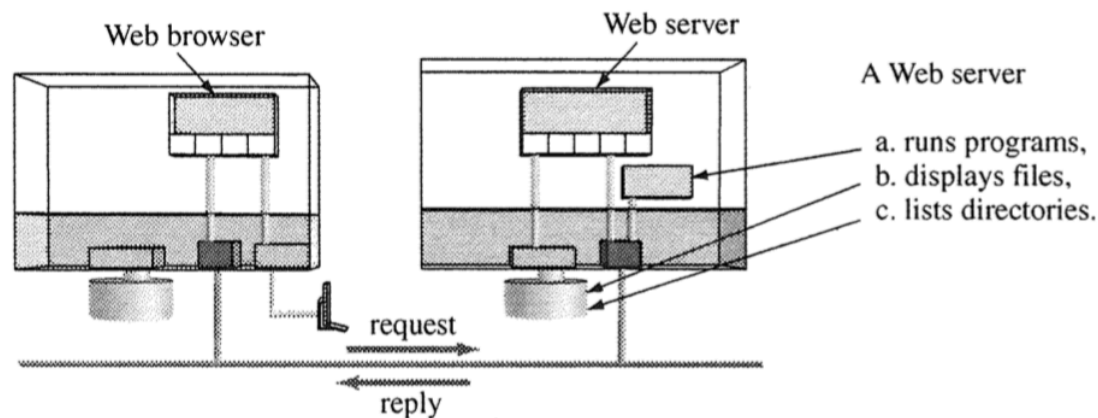
## ■ waitpid()

- First argument is the ID number of the process.
  - -1 tells waitpid to wait for all children.
- Second argument is the pointer to the integer to receive the status.
- The last argument specifies options.
  - WNOHANG option tells waitpid not to wait if there are no zombies.

# What a Web Server Does

---

- A Web server is a simple concept. A Web server is a program that performs the three most common user operations :
  - (a) list directories
  - (b) cat files
  - (c) run programs



# What a Web Server Does

---

## ■ Logic of the Web Server and Client

### client

user selects a link  
connect to server  
write a request

—>

—>

read the reply  
hangup  
display the reply  
    html: render it  
    image: draw it  
    sound: play it  
repeat

<—

### server

accept a call  
read a request  
handle request:  
    directory: list it  
    regular file: cat it  
    .cgi file: run it  
    not exist: error message  
write a reply

# Planning Our Web Server

---

- What operations do we need to code?

## **(a) Set up the server**

We can use `make_server_socket` from `socklib.c`

## **(b) Accept a call**

Use `accept` to get the file descriptor to the client. We can use `fdopen` to make that file descriptor into a buffered stream.

## **(c) Read a request**

What does a request look like? How does the client ask for something? We need to study this one more.

## **(d) handle the request**

We know how to list directories, cat files, and run programs. We can use `opendir` and `readdir`, `open` and `read`, `dup2` and `exec`.

## **(e) Send a reply**

What does a reply look like? What does the client expect to see? This one also requires more study.

# The Protocol of a Web Server

```
$ telnet www.prenhall.com 80
```

```
Trying 165.193.123.253...
```

```
Connected to www.prenhall.com.
```

```
Escape character is '^['.
```

```
GET /index.html HTTP/1.0
```

Connect a web server (port number : 80)

- HTTP request
  - 1<sup>st</sup> Arg: Command
  - 2<sup>nd</sup> Arg: Parameter
  - 3<sup>rd</sup> Arg: version of the protocol

```
HTTP/1.1 200 OK
```

```
Date: Tue, 20 Nov 2018 21:27:27 GMT
```

```
Server: Apache
```

```
X-Powered-By: PHP/5.6.30
```

```
Content-Length: 717
```

```
Connection: close
```

```
Content-Type: text/html; charset=UTF-8
```

HTTP Response

Header

```
<HTML><HEAD>
```

```
<META HTTP-EQUIV="Refresh" CONTENT="0; URL=http://vig.prenhall.com/">
```

```
</HEAD><BODY></BODY></HTML>
```

```
<!-- ----->
```

```
<!--      Caught you peeking!      -->
```

```
<!-- ----->
```

Content

```
Connection closed by foreign host.
```

```
$
```

# Writing a Web Server

---

- The main loop of our Web server looks like the following :

```
while(1) {  
    fd = accept(sock, NULL, NULL);      /* take a call          */  
    fpin = fdopen(fd, "r");              /* make it a FILE *     */  
    fgets(fpin, request, LEN);           /* read client request  */  
    read_until_crnl(fpin);               /* skip over arguments */  
    process_rq(request, fd);             /* reply to client      */  
    fclose(fpin);                       /* hang up connection  */  
}
```



# Webserv Source Code

---

```
/* webserv.c - a minimal web server (version 0.2)
 *
 *  usage: ws portnumber
 *
 *  features: supports the GET command only
 *
 *           runs in the current directory
 *
 *           forks a new child to handle each request
 *
 *           has MAJOR security holes, for demo
purposes only
 *
 *           has many other weaknesses, but is a good
start
 *
 *  build: cc webserv.c socklib.c -o webserv
 */

#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <string.h>
```

```
main(int ac, char *av[])
{
    int          sock, fd;
    FILE         *fpin;
    char         request[BUFSIZ];

    if ( ac == 1 ) {
        fprintf(stderr, "usage: ws portnum\n");
        exit(1);
    }
    sock = make_server_socket( atoi(av[1]) );

    if ( sock == -1 ) exit(2);

    /* main loop here */
    while(1){
        /* take a call and buffer it */
        fd = accept( sock, NULL, NULL );
        fpin = fdopen(fd, "r" );

        /* read request */
        fgets(request, BUFSIZ, fpin);
        printf("got a call: request = %s", request);
        read_til_crnl(fpin);

        /* do what client asks */
        process_rq(request, fd);

        fclose(fpin);
    }
}
```

# Webserv Source Code

---

```
/* ----- */
read_til_crnl(FILE *)
skip over all request info until a CRNL is seen
----- */

read_til_crnl(FILE *fp)
{
    char    buf[BUFSIZ];

    while ( fgets(buf,BUFSIZ,fp) != NULL && strcmp(buf,"\r\n") != 0 )
        ;
}
```

# Webserv Source Code

---

```
process_rq( char *rq, int fd )
```

```
{
```

```
    char  cmd[BUFSIZ], arg[BUFSIZ];
```

```
    /* create a new process and return if not the child */
```

```
    if ( fork() != 0 )
```

```
        return;
```

```
    strcpy(arg, "./");          /* precede args with ./ */
```

```
    if ( sscanf(rq, "%s%s", cmd, arg+2) != 2 )
```

```
        return;
```

```
    if ( strcmp(cmd,"GET") != 0 )
```

```
        cannot_do(fd);
```

```
    else if ( not_exist( arg ) )    do_404(arg, fd );
```

```
    else if ( isadir( arg ) )      do_ls( arg, fd );
```

```
    else if ( ends_in_cgi( arg ) ) do_exec( arg, fd );
```

```
    else                            do_cat( arg, fd );
```

```
}
```

```
/* ----- */
```

```
process_rq( char *rq, int fd )
```

```
do what the request asks for and write reply to fd
```

```
handles request in a new process
```

```
rq is HTTP command: GET /foo/bar.html HTTP/1.0
```

```
----- */
```

# Writing a Web Server

---

```
do_ls(char *dir, int fd)
{
    FILE    *fp;

    fp = fdopen(fd, "w");           /* make socket into a FILE* */
    header(fp, "text/plain");       /* send HTTP reply header */
    fprintf(fp, "\r\n");           /* and end of header mark */
    fflush(fp);                    /* force to socket */

    dup2(fd, 1);                   /* make socket stdout */
    dup2(fd, 2);                   /* make socket stderr */
    close(fd);                     /* close socket */
    execl("/bin/lis", "lis", "-l", dir, NULL); /* lis -l does the work */
    perror(dir);                   /* or it doesn't */
    exit(1);                       /* child exits */
}
```

# Webserv Source Code

---

```
/* ----- *  
the reply header thing: all functions need one  
if content_type is NULL then don't send content type  
----- */  
  
header( FILE *fp, char *content_type )  
{  
    fprintf(fp, "HTTP/1.0 200 OK\r\n");  
    if ( content_type )  
        fprintf(fp, "Content-type: %s\r\n", content_type );  
}
```

# Webserv Source Code

---

```
cannot_do(int fd)
{
    FILE    *fp = fdopen(fd,"w");

    fprintf(fp, "HTTP/1.0 501 Not Implemented\r\n");
    fprintf(fp, "Content-type: text/plain\r\n");
    fprintf(fp, "\r\n");

    fprintf(fp, "That command is not yet implemented\r\n");
    fclose(fp);
}
```

# Webserv Source Code

---

```
do_404(char *item, int fd)
{
    FILE    *fp = fdopen(fd,"w");

    fprintf(fp, "HTTP/1.0 404 Not Found\r\n");
    fprintf(fp, "Content-type: text/plain\r\n");
    fprintf(fp, "\r\n");

    fprintf(fp, "The item you requested: %s\r\nis not found\r\n", item);
    fclose(fp);
}
```

# Webserv Source Code

---

```
/* ----- */
the directory listing section
isadir() uses stat, not_exist() uses stat
do_ls runs ls. It should not
----- */
```

```
isadir(char *f)
{
    struct stat info;
    return ( stat(f, &info) != -1 && S_ISDIR(info.st_mode) );
}
```

```
not_exist(char *f)
{
    struct stat info;
    return( stat(f,&info) == -1 );
}
```

```
do_ls(char *dir, int fd)
{
    FILE    *fp ;

    fp = fdopen(fd,"w");
    header(fp, "text/plain");
    fprintf(fp, "\r\n");
    fflush(fp);

    dup2(fd,1);
    dup2(fd,2);
    close(fd);
    execlp("ls","ls","-l",dir,NULL);
    perror(dir);
    exit(1);
}
```



# Webserv Source Code

---

```
/* ----- *  
the cgi stuff. function to check extension and  
one to run the program.  
----- */
```

```
char * file_type(char *f)
```

```
/* returns 'extension' of file */
```

```
{  
    char *cp;  
    if ( (cp = strrchr(f, '.')) != NULL )  
        return cp+1;  
    return "";  
}
```

```
ends_in_cgi(char *f)
```

```
{  
    return ( strcmp( file_type(f), "cgi" ) == 0 );  
}
```

```
do_exec ( char *prog, int fd )
```

```
{  
    FILE *fp ;  
  
    fp = fdopen(fd,"w");  
    header(fp, NULL);  
    fflush(fp);  
    dup2(fd, 1);  
    dup2(fd, 2);  
    close(fd);  
    execl(prog,prog,NULL);  
    perror(prog);  
}
```

# Webserv Source Code

---

```
/*-----*  
do_cat(filename,fd)  
sends back contents after a header  
-----*/
```

```
do_cat(char *f, int fd)
```

```
{  
    char    *extension = file_type(f);  
    char    *content = "text/plain";  
    FILE    *fpsock, *fpfile;  
    int     c;  
  
    if ( strcmp(extension,"html") == 0 )  
        content = "text/html";  
    else if ( strcmp(extension, "gif") == 0 )  
        content = "image/gif";  
    else if ( strcmp(extension, "jpg") == 0 )  
        content = "image/jpeg";  
    else if ( strcmp(extension, "jpeg") == 0 )  
        content = "image/jpeg";
```

```
    fpsock = fdopen(fd, "w");  
    fpfile = fopen( f , "r");
```

```
    if ( fpsock != NULL && fpfile != NULL )  
    {  
        header( fpsock, content );  
        fprintf(fpsock, "\r\n");  
  
        while( (c = getc(fpfile) ) != EOF )  
            putc(c, fpsock);  
  
        fclose(fpfile);  
        fclose(fpsock);  
    }  
    exit(0);  
}
```

# Running the Web Server

---

- Compile the code and then run the program at a port:
  - `$ cc webserv.c socklib.c -o webserv`
  - `$ ./webserv 12345`
- Create the following shell script (hello.cgi):

```
#!/bin/sh
# hello.cgi - a cheery cgi page
printf "Content-type: text/plain\n\nhello\n";
```

- call this script hello.cgi, chmod it to 755,
- and then your browser to invoke the script with
  - `http://yourhostname:12345/`
  - `http://yourhostname:12345/hello.cgi`
  - `http://yourhostname:12345/webserv.c`
  - `http://yourhostname:12345/mypage`